


AI61003 LINEAR ALGEBRA FOR AI AND ML
ASSIGNMENT 01 - PROBLEM 09

N (the number of training sample) : **1000**

n (the dimension of vectors to be clustered) : **784**

CASE 1 Random Initialization of Cluster Representatives



Number of iterations taken for the algorithm to converge : **18**


Accuracy of cluster assignment for a randomly selected set of 50 images : **68.00%**

| Number of clusters (k) | J_{clust} |
|------------------------|--------------------|
| 5 | 5.8968 |
| 6 | 5.8798 |
| 7 | 5.8824 |
| 8 | 5.8515 |
| 9 | 5.8549 |
| 10 | 5.8671 |
| 11 | 5.8752 |
| 12 | 5.9129 |
| 13 | 5.8592 |

| | |
|----|--------|
| 14 | 5.8509 |
| 15 | 5.8954 |
| 16 | 5.8937 |
| 17 | 5.8979 |
| 18 | 5.8635 |
| 19 | 5.8869 |
| 20 | 5.8572 |

The optimal number of clusters is the one that gives the least value of the objective function J_{clust} . Therefore, **14** is the optimal cluster size.

CASE 2 Choosing Cluster Representatives from Given Dataset



Number of iterations taken for the algorithm to converge : **16**
 Accuracy of cluster assignment for a randomly selected set of 50 images : **72.00%**

| Number of clusters (k) | J_{clust} |
|------------------------|--------------------|
| 5 | 6.5726 |
| 6 | 6.4847 |
| 7 | 6.4099 |
| 8 | 6.3861 |

| | |
|----|--------|
| 9 | 6.3043 |
| 10 | 6.2270 |
| 11 | 6.1661 |
| 12 | 6.1165 |
| 13 | 6.1421 |
| 14 | 6.0424 |
| 15 | 6.0164 |
| 16 | 5.9788 |
| 17 | 5.9403 |
| 18 | 5.9260 |
| 19 | 5.9095 |
| 20 | 5.8953 |

The optimal number of clusters is the one that gives the least value of the objective function J_{clust} . Therefore, **20** is the optimal cluster size.

The choice of initial condition definitely has an effect on the performance of the clustering algorithm.

- Firstly the algorithm converges faster in the second case than in the first case. The converged value of inertia (or *within-cluster sum-of-squares criterion*) in the first case (for $k = 20$) was 35668 and in the second case was 35658. It shows that not only does the number of iterations get affected by the initial condition, but in the second case the algorithm may tend to find a more optimal clustering (converges towards global minima) than in the first case (might converge towards local minima). Having said this, though both the number of iterations and the converged criterion value are better for *data-dependent initialization* of cluster representatives, the relative difference between these values for the two cases is quite small. So, random initialization of cluster representatives can also give a good clustering of objects.
- The accuracy of cluster assignments was evaluated in both the cases on the same test set (constructed by randomly choosing 50 images from MNIST dataset that were not seen in training). It is clear that the accuracy of cluster assignment in the second case is 4% higher than in the first case. This shows that the initial condition can also have an impact on the performance of the produced clusters on the unseen test samples.
- If you analyze the change in the converging value of J_{clust} with the number of clusters, you would find that the converging value decreases (almost) monotonically in the second case, whereas in the first case there is a very irregular pattern. As a result of this, their optimal cluster sizes also come out to be very different (14 for random initialization and 20 for data-dependent initialization).

```
In [1]: 1 import tensorflow as tf
        2 import numpy as np
        3 from matplotlib import pyplot as plt
        4 from random import shuffle, randint
        5 import sklearn.cluster as Clustering_Algorithm
```

```
In [2]: 1 mnist_data = tf.keras.datasets.mnist
        2 (x_train, y_train), (x_test, y_test) = mnist_data.load_data()
        3 train_data = list(zip(x_train, y_train))
        4 shuffle(train_data)
        5 x_train, y_train = list(zip(*train_data))
```

```
In [3]: 1 train_images = [ [] for digit in range(10) ]
        2 sample_count_per_digit = 100
        3 for i, label in enumerate(y_train) :
        4     if ( len(train_images[label]) == sample_count_per_digit ) : continue
        5     train_images[label].append(x_train[i].flatten().astype(np.float32) / 255)
        6
        7 def Get_Random_Test_Data ( test_samples_count = 50 ) :
        8     global x_test, y_test
        9
        10    test_data = list(zip(x_test, y_test))
        11    shuffle(train_data)
        12    shuffle(test_data)
        13    x_train, y_train = list(zip(*train_data))
        14    x_test_shuff, y_test_shuff = list(zip(*test_data))
        15
        16    test_images = [ [] for digit in range(10) ]
        17    count = 0
        18    for i, label in enumerate(y_test_shuff) :
        19        if count == test_samples_count : break
        20        count += 1
        21        test_samples_count
        22        test_images[label].append(x_test_shuff[i].flatten().astype(np.float32) / 255)
        23    return test_images
```

```
In [4]: 1 train = []
        2 do = [train.extend(img) for img in train_images]
        3 N = len(train)
        4 n = train_images[0][0].shape[0]
        5 print ( 'N (number of training samples) :', N)
        6 print('n (dimension of vectors) :', n)
```

N (number of training samples) : 1000

n (dimension of vectors) : 784

```
In [5]: 1 def Get_Distance ( a , b ) :
        2     return np.linalg.norm(b-a)
```

```
In [6]: 1 def J_clust ( train , kmeans ) :
        2     s = 0
        3     for img in train :
        4         cluster_id = kmeans.predict(img.astype(np.float).reshape(1,-1))[0]
        5         s += Get_Distance(img, kmeans.cluster_centers_[cluster_id])
        6     return s / len(train)
```

CASE I : Random Initialization of Cluster Representatives

In [7]:

```
1 k = 20
2 kmeans = Clustering_Algorithm.KMeans(n_clusters = k, init='random', verbose=True, n_init=1)
3 kmeans.fit(train)
```

Initialization complete

Iteration 0, inertia 60732.524707579985

Iteration 1, inertia 38050.91288994666

Iteration 2, inertia 36886.65874578197

Iteration 3, inertia 36505.343118868186

Iteration 4, inertia 36233.53763940629

Iteration 5, inertia 36055.34772788261

Iteration 6, inertia 35925.409068793735

Iteration 7, inertia 35836.56467969832

Iteration 8, inertia 35781.09028613918

Iteration 9, inertia 35733.11039512983

Iteration 10, inertia 35708.260845345685

Iteration 11, inertia 35688.57368506976

Iteration 12, inertia 35678.89407890276

Iteration 13, inertia 35675.8489286663

Iteration 14, inertia 35672.49733949574

Iteration 15, inertia 35671.12267063208


Iteration 16, inertia 35669.257571344424

Iteration 17, inertia 35668.07054899751

Converged at iteration 17: strict convergence.

Out[7]: KMeans(init='random', n_clusters=20, n_init=1, verbose=True)

```
In [8]: 1 for cluster_id, rep in enumerate(kmeans.cluster_centers_) :
2         plt.subplot(4, 5, cluster_id+1)
3         cluster_rep_img = (rep*255).astype(np.int32).reshape((28,28))
4         plt.imshow(cluster_rep_img, cmap='gray')
5         plt.title('z_' + str(cluster_id+1), fontsize=9, pad=4)
6         plt.axis('off')
7 plt.savefig('cluster_representative_plots_part_1.png')
```



```
In [9]: 1 digits_represented_by_clusters = [5, 1, 6, 4, 9, 2, 7, 2, 1, 5, 3, 8, 0, 9, 6, 6, 0, 9, 8, 0]
```

```
In [10]: 1 correctly_classified = 0
2 test_images_count = 50
3 test_images = Get_Random_Test_Data(test_images_count)
4 for dig in range(10) :
5     for img in test_images[dig] :
6         cluster_id = kmeans.predict(img.astype(np.float).reshape(1,-1))[0]
7         correctly_classified += (digits_represented_by_clusters[cluster_id] == dig)
8
9 accuracy = 100 * correctly_classified / test_images_count
10 print('ACCURACY OF CLUSTER ASSIGNMENT :', round(accuracy, 3))
```

ACCURACY OF CLUSTER ASSIGNMENT : 68.0

```
In [11]: 1 least_J_value = float('inf')
2 opt_cluster_size = -1
3 for cluster_size in range(5, 21) :
4     kmeans = Clustering_Algorithm.KMeans(n_clusters = k, init='random', verbose=False, n_init=1)
5     kmeans.fit(train)
6     J_value = J_clust(train, kmeans)
7     print('CLUSTER SIZE :', cluster_size, '\tJ_clust :', round(J_value, 5))
8
9     if J_value < least_J_value :
10         least_J_value = J_value
11         opt_cluster_size = cluster_size
12
13 print('OPTIMAL CLUSTER SIZE :', opt_cluster_size)
```

```
CLUSTER SIZE : 5          J_clust : 5.89679
CLUSTER SIZE : 6          J_clust : 5.87985
CLUSTER SIZE : 7          J_clust : 5.8824
CLUSTER SIZE : 8          J_clust : 5.85147
CLUSTER SIZE : 9          J_clust : 5.8549
CLUSTER SIZE : 10         J_clust : 5.86706
CLUSTER SIZE : 11         J_clust : 5.8752
CLUSTER SIZE : 12         J_clust : 5.91291
CLUSTER SIZE : 13         J_clust : 5.85923
CLUSTER SIZE : 14         J_clust : 5.85086
CLUSTER SIZE : 15         J_clust : 5.89545
CLUSTER SIZE : 16         J_clust : 5.8937
CLUSTER SIZE : 17         J_clust : 5.89795
CLUSTER SIZE : 18         J_clust : 5.86347
CLUSTER SIZE : 19         J_clust : 5.88693
CLUSTER SIZE : 20         J_clust : 5.85719
OPTIMAL CLUSTER SIZE : 14
```

CASE II : Data Dependent Initialization of Cluster Representatives



```
In [12]: 1 def Data_Dependent_Cluster_Representatives_Init ( train , k ) :
2         clusters = []
3         for i in range(k) : clusters.append([])
4         for img in train :
5             clusters[randint(0, k-1)].append(img)
6
7         cluster_representatives = []
8         for i in range(k) :
9             rep = sum(clusters[i]) / len(clusters[i])
10            cluster_representatives.append(rep)
11
12        cluster_representatives = np.array(cluster_representatives)
13        return cluster_representatives
```

```
In [13]: 1 k = 20
2 kmeans = Clustering_Algorithm.KMeans(n_clusters = k, verbose=True, n_init=1,
3                                     init=Data_Dependent_Cluster_Representatives_Init(train, k))
4 kmeans.fit(train)
```

```
Initialization complete
Iteration 0, inertia 49601.48328326748
Iteration 1, inertia 39424.78061696993
Iteration 2, inertia 37262.66951403245
Iteration 3, inertia 36526.05086918052
Iteration 4, inertia 36129.293411810526
Iteration 5, inertia 35995.34874088044
Iteration 6, inertia 35928.17828729691
Iteration 7, inertia 35885.53611488616
Iteration 8, inertia 35829.41673297419
Iteration 9, inertia 35792.76307584839
Iteration 10, inertia 35752.53303179572
Iteration 11, inertia 35721.684346373004
Iteration 12, inertia 35687.63237248872
Iteration 13, inertia 35672.999806197186
Iteration 14, inertia 35664.05230663018
Iteration 15, inertia 35658.41160480236
Converged at iteration 15: strict convergence.
```

```
Out[13]: KMeans(init=array([[0., 0., 0., ..., 0., 0., 0.],
    [0., 0., 0., ..., 0., 0., 0.],
    [0., 0., 0., ..., 0., 0., 0.],
    ...,
    [0., 0., 0., ..., 0., 0., 0.],
    [0., 0., 0., ..., 0., 0., 0.],
    [0., 0., 0., ..., 0., 0., 0.]]), dtype=float32),
    n_clusters=20, n_init=1, verbose=True)
```

```
In [15]: 1 for cluster_id, rep in enumerate(kmeans.cluster_centers_) :
2         plt.subplot(4, 5, cluster_id+1)
3         cluster_rep_img = (rep*255).astype(np.int32).reshape((28,28))
4         plt.imshow(cluster_rep_img, cmap='gray')
5         plt.title('z_' + str(cluster_id+1), fontsize=9, pad=4)
6         plt.axis('off')
7 plt.savefig('cluster_representative_plots_part_2.png')
```



```
In [16]: 1 digits_represented_by_clusters = [0, 5, 6, 9, 6, 9, 6, 9, 8, 2, 1, 7, 1, 1, 7, 5, 3, 0, 9, 9]
```

```
In [17]: 1 correctly_classified = 0
2 test_images_count = 50
3 for dig in range(10) :
4     for img in test_images[dig] :
5         cluster_id = kmeans.predict(img.astype(np.float).reshape(1,-1))[0]
6         correctly_classified += (digits_represented_by_clusters[cluster_id] == dig)
7
8 accuracy = 100 * correctly_classified / test_images_count
9 print('ACCURACY OF CLUSTER ASSIGNMENT :', round(accuracy, 3))
```

ACCURACY OF CLUSTER ASSIGNMENT : 72.0

```
In [18]: 1 least_J_value = float('inf')
2 opt_cluster_size = -1
3 for cluster_size in range(5, 21) :
4     kmeans = Clustering_Algorithm.KMeans(n_clusters = cluster_size, verbose=False, n_init=1,
5                                           init=Data_Dependent_Cluster_Representatives_Init(train, cluster_size))
6     kmeans.fit(train)
7     J_value = J_clust(train, kmeans)
8     print('CLUSTER SIZE :', cluster_size, '\tJ_clust :', round(J_value, 5))
9
10    if J_value < least_J_value :
11        least_J_value = J_value
12        opt_cluster_size = cluster_size
13
14 print('OPTIMAL CLUSTER SIZE :', opt_cluster_size)
```

```
CLUSTER SIZE : 5          J_clust : 6.57261
CLUSTER SIZE : 6          J_clust : 6.48468
CLUSTER SIZE : 7          J_clust : 6.40986
CLUSTER SIZE : 8          J_clust : 6.38606
CLUSTER SIZE : 9          J_clust : 6.30426
CLUSTER SIZE : 10         J_clust : 6.22705
CLUSTER SIZE : 11         J_clust : 6.16606
CLUSTER SIZE : 12         J_clust : 6.11653
CLUSTER SIZE : 13         J_clust : 6.14211
CLUSTER SIZE : 14         J_clust : 6.04239
CLUSTER SIZE : 15         J_clust : 6.0164
CLUSTER SIZE : 16         J_clust : 5.97875
CLUSTER SIZE : 17         J_clust : 5.94029
CLUSTER SIZE : 18         J_clust : 5.92599
CLUSTER SIZE : 19         J_clust : 5.90952
CLUSTER SIZE : 20         J_clust : 5.8953
OPTIMAL CLUSTER SIZE : 20
```