

```
import torch
import torch.nn as nn
import json
import pickle
import numpy as np
from random import shuffle, random
from time import time
import seaborn as sns
import matplotlib.pyplot as plt
from collections import Counter
from math import log2
```

```
In [2]: # Extract data from the specified location
def Extract_Data ( loc ) :
    file = open(loc, 'r')
    data = file.read()
    data = data.split('\n')[1:-1]
    data = [ x.split(',') for x in data ]
    data = [ (x[i-1], eval(x[i-1])) for x in data ]
    data = [ (eval(x) for x in X, _ ) for X, _ in data ]
    return data

# Given the probability distribution of a discrete attribute for its different values
# in the form of a vector, find the entropy of that attribute.
def Entropy ( dist ) :
    vals = np.array(dist.values())
    vals = vals / vals.sum()
    s = sum(vals * np.log2(vals))
    if (s == 0.0) : return 0
    return -1 * s

# Given the data matrix, study the various properties of the attributes.
def Study_Attributes ( data_matrix ) :
    std_dev_atts = np.std(data_matrix, axis = 0)

    # Study standard deviation in the values of the attributes
    fig = plt.figure()
    ax = fig.add_subplot()
    sns.set(font_scale = 1.2)
    sns.heatmap(std_dev_atts.reshape((8, 8)))
    plt.title('STD. DEVIATION IN THE ATTRIBUTE VALUES\n', fontsize=13, fontweight='bold')
    ax.axes.get_xaxis().set_visible(False)
    ax.axes.get_yaxis().set_visible(False)
    plt.plot()
    plt.savefig('PREPROCESSING_std_dev_attributes.png')

    # Study diversity in the values of the attributes
    distinct_att_vals = np.zeros(64, dtype = np.int32)
    attribute_entropies = [ ]
    for att in range(64) :
        dist = Counter(data_matrix[:, att])
        distinct_att_vals[att] = len(dist.keys())
        attribute_entropies.append(Entropy(dist))
    attribute_entropies = np.array(attribute_entropies)

    fig = plt.figure()
    ax = fig.add_subplot()
    sns.set(font_scale = 1.2)
    sns.heatmap(distinct_att_vals.reshape((8, 8)))
    plt.title('DIVERSITY IN THE ATTRIBUTE VALUES\n', fontsize=13, fontweight='bold')
    ax.axes.get_xaxis().set_visible(False)
    ax.axes.get_yaxis().set_visible(False)
    plt.plot()
    plt.savefig('PREPROCESSING_diversity_attributes.png')

    # Study entropy in the values of the attributes
    fig = plt.figure()
    ax = fig.add_subplot()
    sns.set(font_scale = 1.2)
    sns.heatmap(attribute_entropies.reshape((8, 8)))
    plt.title('ENTROPY OF THE ATTRIBUTES\n', fontsize=13, fontweight='bold')
    ax.axes.get_xaxis().set_visible(False)
    ax.axes.get_yaxis().set_visible(False)
    plt.plot()
    plt.savefig('PREPROCESSING_entropy_attributes.png')

    return std_dev_atts, distinct_att_vals, attribute_entropies

# Filter those attributes that have entropy below some threshold. Please refer
# to the report for more theoretical details!
def Filter_Attributes ( data , entropy_thresh = 0.1 ) :
    X = np.array([ X for X, _ in data ])
    _, attribute_entropies = Study_Attributes(X)
    data_filtered = [t for t in range(64) if attribute_entropies[t] >= entropy_thresh], y) for X, y in data
    return data_filtered
```

```
In [3]: # Load data from the given locations. One location is for
# the training data and the other is for test data.
def Load_Data ( loc1, loc2 ) :
    data = Extract_Data(loc1) + Extract_Data(loc2) # merge the sub-datasets
    data = Filter_Attributes(data) # filter attributes
    return data

# Split the entire dataset into train & test in the given ratio.
def Split ( data , train_split_ratio = 0.8 ) :
    shuffle(data)
    size = round(train_split_ratio * len(data))
    train_set = data[:size]
    test_set = data[size:]
    return train_set, test_set

# Train & test both should have a balanced frequency of data samples
# from each of the classes otherwise unreliable results can be obtained.
class samples = [ (xy for xy in data if xy[i] == c) for c in range(10) ]
[ shuffle(_) for _ in class_samples ]

train_set = [ ]
test_set = [ ]
for c in range(10) :
    b = random() * 0.4 - 0.2 # small deviation
    s = round((train_split_ratio + b) * len(class_samples[c]))
    train_set.extend(class_samples[c][:s])
    test_set.extend(class_samples[c][s:])

shuffle(train_set)
shuffle(test_set)
return train_set, test_set
```

```
In [4]: # Package the train data into mini-batches of the given size and
# return the list of batches.
def Make_Batches ( train_data , batch_size = 32 ) :
    shuffle(train_data)
    batches = list()
    for start in range(0, len(train_data), batch_size) :
        end = start + batch_size
        batch_labels = list(zip(*train_data[start:end]))
        labels = torch.Tensor(labels).long()
        batch = torch.Tensor(batch).float()
        batches.append((batch, labels))
    return batches
```

```
In [5]: # Multi-Layer Perceptron Classifier
class Handwritten_Digits_Classifier ( nn.Module ) :
    def __init__ ( self , name , input_size , hidden_layer_sizes ,
                  output_size , lr = 0.01 , momentum = 0.8 , early_stopping_tol = 4 ) :
        super(Handwritten_Digits_Classifier, self).__init__()
        self.__name__ = name
        self.__input_size__ = input_size
        self.__hidden_layer_sizes__ = hidden_layer_sizes
        self.__output_size__ = output_size

        self.__hidden_layers__ = [ ]
        hidden_input_size = input_size
        for size in hidden_layer_sizes :
            layer = nn.Linear(hidden_input_size, size, bias = True)
            self.__hidden_layers__.append(layer)
            hidden_input_size = size

        self.__output_layer__ = nn.Linear(hidden_input_size, output_size, bias = True)
        self.__loss__ = nn.CrossEntropyLoss()
        self.__optimizer__ = torch.optim.SGD(self.parameters(), lr = lr, momentum = momentum)

        self.__training_history__ = {
            'train_loss_progress' : [ ],
            'train_acc_progress' : [ ],
            'valid_acc_progress' : [ ]
        }

        # Additional information
        self.__stopping_reason__ = 'Kernel interrupted'
        self.__learning_rate__ = lr
        self.__momentum__ = momentum
        self.__early_stopping_tol__ = early_stopping_tol
        self.__training_time__ = None
        self.__batch_size__ = None

    def forward ( self , batch_input ) :
        batch_output = batch_input
        for layer in self.__hidden_layers__ :
            batch_output = layer(batch_output)
        batch_output = self.__output_layer__(batch_output)
        return batch_output

    def backward ( self , pred_label_dist , true_labels ) :
        loss = self.__loss__(pred_label_dist, true_labels)
        self.__optimizer.zero_grad()
        loss.backward()
        self.__optimizer.step()
        return loss.item()

    def accuracy ( self , data ) :
        self.eval()
        total_correct = 0
        tensor_samples = Make_Batches(data, 1)
        for _, (sample, label) in enumerate(tensor_samples) :
            label_dist = self(sample)
            pred = torch.argmax(label_dist, dim = 1)
            corr_count = (pred == label).sum()
            total_correct += corr_count
        self.train()
        return total_correct.item() / len(data)

    def loss ( self , data ) :
        self.eval()
        total_loss = 0
        tensor_samples = Make_Batches(data, 1)
        for _, (sample, label) in enumerate(tensor_samples) :
            label_dist = self(sample)
            total_loss += self.__loss__(label_dist, label).item()
        self.train()
        return total_loss / len(data)

    def predict ( self , x ) :
        X = np.array(x, dtype = np.float32).reshape((1, -1))
        X = torch.from_numpy(X)
        label_dist = self(X)
        pred = torch.argmax(label_dist, dim = 1)
        return pred.item()

    def confusion_matrix ( self , data ) :
        cf = np.zeros((10, 10), dtype = np.int32)
        for x, y in data :
            cf[y][self.predict(x)] += 1
        return cf

    def plot_confusion_matrix ( self , data , loc , title = 'CONFUSION MATRIX (TRAIN DATA)' ) :
        cf = self.confusion_matrix(data)
        fig = plt.figure(figsize=(4, 3))
        ax = fig.add_subplot()
        sns.set(font_scale = 1.0)
        sns.heatmap(cf, np.sum(cf, axis = 1).reshape(10, 1))
        plt.title(self.__name__ + '\n' + title, fontsize=11, fontweight='bold')
        plt.xlabel('PREDICTED CLASS', fontsize=9)
        plt.ylabel('TRUE CLASS', fontsize=9)
        plt.tight_layout()
        plt.show()

    def plot_training_history ( self , loc ) :
        fontsize = 12
        plt.style.use('seaborn-whitegrid')
        epochs = len(self.__training_history['train_loss_progress'])
        X = list(range(1, 1 + epochs))

        plt.figure(figsize=(4.2, 3.2))
        plt.plot(X, self.__training_history['train_acc_progress'], linewidth = 2.2)
        plt.plot(X, self.__training_history['valid_acc_progress'], linewidth = 1.9)
        plt.xlabel('EPOCH', fontsize=fontsize)
        plt.ylabel('ACCURACY', fontsize=fontsize)
        plt.legend(['Train Acc.', 'Validation Acc.'], fontsize = 9)
        plt.title(self.__name__ + ' -- ACCURACY PROGRESS', fontsize=11, fontweight='bold')
        plt.tight_layout()
        plt.show()

        plt.figure(figsize=(4.2, 3.2))
        plt.plot(X, self.__training_history['train_loss_progress'], linewidth = 2, color = '#d62728')
        plt.xlabel('EPOCH', fontsize=fontsize)
        plt.ylabel('LOSS', fontsize=fontsize)
        plt.title(self.__name__ + ' -- LOSS PROGRESS', fontsize=11, fontweight='bold')
        plt.tight_layout()
        plt.show()

        # Print a detailed report about training, accuracies, losses etc including the
        # plots and confusion matrices.
        def model_analysis_report ( self , train_data , test_data ) :
            Print('+++ MODEL ANALYSIS REPORT +++')

            print('\n >> INTRODUCTION')
            print(' * NAME :', self.__name__)
            print(' * INPUT SIZE :', self.__input_size__)
            print(' * OUTPUT SIZE :', self.__output_size__)

            print('\n >> HYPERPARAMETERS')
            print(' * HIDDEN LAYERS SIZES :', self.__hidden_layer_sizes__)
            print(' * LEARNING RATE :', self.__learning_rate__)
            print(' * BATCH SIZE :', self.__batch_size__)
            print(' * MOMENTUM :', self.__momentum__)
            print(' * EARLY STOPPING TOLERANCE :', self.__early_stopping_tol__)

            print('\n >> TRAINING INFORMATION')
            print(' * EPOCHS :', len(self.__training_history['train_loss_progress']))
            print(' * TRAINING TIME : {:.3f} mins'.format(self.__training_time / 60))
            print(' * STOPPING REASON :', self.__stopping_reason__)
            print(' * STOPPING TRAINING LOSS : {:.5f}'.format(self.__training_history['train_loss_progress'][-1]))
            print(' * STOPPING TRAINING ACCURACY : {:.3f} %'.format(100 * self.__training_history['train_acc_progress'][-1]))
            print(' * STOPPING VALIDATION ACCURACY : {:.3f} %'.format(100 * self.__training_history['valid_acc_progress'][-1]))

            print('\n >> TRAINING PROGRESS')
            self.__plot_training_history(self.__name__ + str(self.__learning_rate)[2:] + 'prog.png')

            print('\n >> MODEL TESTING - ACCURACY & LOSS')
            train_acc = 100 * self.accuracy(train_data)
            test_acc = 100 * self.accuracy(test_data)
            print(' * TRAIN ACCURACY : {:.3f} %'.format(train_acc))
            print(' * TEST ACCURACY : {:.3f} %'.format(test_acc))
            print(' * TRAIN LOSS : {:.5f}'.format(self.__loss__(train_data)))
            print(' * TEST LOSS : {:.5f}'.format(self.__loss__(test_data)))

            print('\n >> MODEL TESTING - CONFUSION MATRICES')
            self.__plot_confusion_matrix(train_data, self.__name__ + str(self.__learning_rate)[2:] + 'train_cf.png',
                                       title = 'CONFUSION MATRIX (TRAIN DATA)')
            self.__plot_confusion_matrix(test_data, self.__name__ + str(self.__learning_rate)[2:] + 'test_cf.png',
                                       title = 'CONFUSION MATRIX (TEST DATA)')

            return train_acc, test_acc

        # Strict training is used when we have to train a model for a specific number of epochs, no matter if
        # the model starts to over-fit or the train loss starts increasing. So the model must be trained exactly
        # for a certain number of epochs, without early-stopping, therefore there is no need for validation set.
        # Though, the accuracy of the model on the test set is registered while training, after every epoch. The
        # type of training is used only while tuning the hyperparameter "Number of Training Epochs" in the 4th param
        def strict_train ( self , total_epochs , train_data , test_data ,
                          batch_size = 16 , verbose = False , verboseGap = 1 ) :
            self.train()
            train_acc_prog = [ ]
            test_acc_prog = [ ]
            for epoch in range(1, 1 + total_epochs) :
                total_loss = 0.0
                total_correct = 0
                batches = Make_Batches(train_data, batch_size)
                for _, (batch, labels) in enumerate(batches) :
                    label_dist = self(batch)
                    loss = self.backward(label_dist, labels)
                    total_loss += loss
                    pred = torch.argmax(label_dist, dim = 1)
                    corr_count = (pred == labels).sum()
                    total_correct += corr_count

                train_acc = self.accuracy(train_data)
                avg_loss = round(self.__loss__(train_data), 5)
                test_acc = round(self.accuracy(test_data), 5)
                train_acc_prog.append(train_acc)
                test_acc_prog.append(test_acc)

                if verbose and (epoch - 1) % verboseGap == 0 :
                    print('EPOCH :{:3d} | TRAIN LOSS : {:.5f} | TRAIN ACC. : {:.5f} | TEST ACC. : {:.5f} '.format(epoch, avg_loss, train_acc, test_acc))

            self.eval()
            return train_acc_prog, test_acc_prog

        # Training on the train data. The training can terminate due to -
        # EARLY STOPPING (very common) - when validation accuracy starts decreasing or train loss starts increasing
        # MAX EPOCHS REACHED (rare) - when the model is trained for the maximum allowed number of epochs.
        # TERMINAL ACCURACY REACHED (very rare) - when model's accuracy on train set exceeds a very high threshold.
        def train ( self , train_data , validation_data , validation_ratio = 0.2 , batch_size = 16 ,
                  total_epochs = 200 , terminal_train_acc = 99.0 , verbose = False ) :
            self.__batch_size__ = batch_size
            self.train()
            self.__name__ = (self.__name__ + str(batch_size))
            L = len(train_data)

            train_size = int(0.8 * L)
            train_set = train_data[:train_size]
            valid_set = train_data[train_size:]

            train_loss_change = [ 1 ] * self.__early_stopping_tol
            valid_acc_change = [ 1 ] * self.__early_stopping_tol
            last_train_loss = float('inf')
            last_valid_acc = 0.0
            start_time = time()

            for epoch in range(1, 1 + total_epochs) :
                total_loss = 0.0
                total_correct = 0
                batches = Make_Batches(train_set, batch_size)
                for _, (batch, labels) in enumerate(batches) :
                    label_dist = self(batch)
                    loss = self.backward(label_dist, labels)
                    total_loss += loss
                    pred = torch.argmax(label_dist, dim = 1)
                    corr_count = (pred == labels).sum()
                    total_correct += corr_count

                train_acc = self.accuracy(train_set)
                avg_loss = round(self.__loss__(train_set), 5)
                valid_acc = round(self.accuracy(valid_set), 5)

                self.__training_history['train_loss_progress'].append(avg_loss)
                self.__training_history['train_acc_progress'].append(train_acc)
                self.__training_history['valid_acc_progress'].append(valid_acc)

                if verbose :
                    print('EPOCH :{:3d} | TRAIN LOSS : {:.5f} | TRAIN ACC. : {:.5f} | VALIDATION ACC. : {:.5f} '.format(epoch, avg_loss, train_acc, valid_acc))

                if train_acc >= terminal_train_acc :
                    self.__training_time = time() - start_time
                    if verbose : print('\n ( STOPPING : Terminal training accuracy reached )')
                    self.__stopping_reason__ = 'Terminal training accuracy was reached'
                    self.eval()
                    return

                if valid_acc >= last_valid_acc :
                    valid_acc_change = valid_acc_change[1:] + [ 1 ]
                    last_valid_acc = valid_acc
                else :
                    valid_acc_change = valid_acc_change[1:] + [ 0 ]
                    if sum(valid_acc_change) == 0 :
                        self.__training_time = time() - start_time
                        self.__stopping_reason__ = [ 'EARLY STOPPING : Validation Accuracy not increasing' ]
                        self.eval()
                        return
                    last_valid_acc = valid_acc

                if avg_loss <= last_train_loss :
                    train_loss_change = train_loss_change[1:] + [ 1 ]
                    last_train_loss = avg_loss
                else :
                    train_loss_change = train_loss_change[1:] + [ 0 ]
                    if sum(train_loss_change) == 0 :
                        self.__training_time = time() - start_time
                        if verbose : print('\n ( STOPPING : Training Loss not decreasing )')
                        self.__stopping_reason__ = [ 'EARLY STOPPING : Training loss was not decreasing' ]
                        self.eval()
                        return
                    last_train_loss = avg_loss

                self.__training_time = time() - start_time
                if verbose : print('\n ( STOPPING : Max no. of training epochs reached )')
                self.__stopping_reason__ = 'Maximum number of epochs was reached'
                self.eval()
                return
```

```
In [6]: def Data_Matrix ( data ) :
        X = [ x for x, _ in data ]
        X = np.array(X)
        X = X.transpose()
        return X

    def Covariance_Matrix ( data ) :
        L = len(data)
        X = Data_Matrix(data)
        mean = X.mean(axis = 1)
        mean = np.tile(mean, (N,1)).transpose()
        X_minus_mean = X - mean
        cov = np.matmul(X_minus_mean, X_minus_mean.transpose())
        return cov / N

# Return the first and second principal component directions for the
# given data matrix.
def PCA_Component_Weights ( data , verbose = False ) :
    cov = Covariance_Matrix(data)
    d = cov.shape[0]
    if verbose :
        eig = list(np.linalg.eig(cov)[0])
        eig.sort(reverse = True)
        for i in range(len(eig)) :
            print(' EIGEN VAL :{:2d} : {:.5f}'.format(i+1, eig[i]))
        eig = np.linalg.eig(cov)
        W = [ i, eig[i][1]] for i in range(d) ] # all the components
        W.sort(key = lambda k: eig[0][k][0], reverse = True)
        # sort the components by their corresponding eigen-value
        W = [ t for _, t in W ]
        return np.array(W[:2]) # return the 1st & 2nd principal component directions

# Given the matrix W, this function returns the reduced dimensionality
# data points for the given data (reduced to 2 dimensions).
def Two_Dim_Points ( data , W ) :
    x_higher = Data_Matrix(data)
    x_2d = np.matmul(W, x_higher).transpose()
    L = len(data)
    transformed = [ [list(x_2d[i]), data[i][1]] for i in range(L) ]
    return transformed
```

```
In [7]: # Plots 2d samples (reduced to 2 dimensions) on a Cartesian plane, belonging to the given classes
def Plot_2D_Samples ( data_full , classes = list(range(10)) , s = 3 , alpha = 1 ) :
    data = [ [x, y] for x, y in data_full if y in classes ]
    points = [ p for p, _ in data ]
    x, y = list(zip(*points))
    colors = ['f8a3ff', '#ff8c00', 'tab:green', '#fffa', '#d2d4ff',
              '#bf8040', '#ff4da6', '#dcdcdc', '#fffffa', '#8cffff']
    col = [colors[t[1]] for t in data]

    plt.figure(figsize=(4, 3))
    plt.scatter(x, y, color = col, s = s, alpha = alpha)
    plt.grid(False)

    filename = 'PART5_classes'
    for c in classes :
        filename += '-' + str(c)
    filename += '.png'

    title = 'CLASSES - ' + str(classes)
    plt.title(title, fontsize=11, fontweight='bold')
    plt.tight_layout()
    plt.savefig(filename)
    plt.show()
    plt.style.use('default')
```

```
In [8]: def Plot_All_Classes_In_2D ( data_full ) :
        colors = ['f8a3ff', '#ff8c00', 'tab:green', '#fffa', '#d2d4ff',
                  '#bf8040', '#ff4da6', '#dcdcdc', '#fffffa', '#8cffff']

        for c in range(10) :
            data = [ [x, y] for x, y in data_full if y == c ]
            points = [ p for p, _ in data ]
            x, y = list(zip(*points))
            col = [colors[t[1]] for t in data]

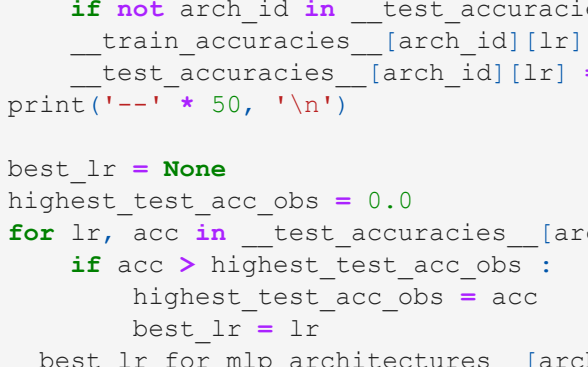
            plt.style.use('dark_background')
            plt.grid(False)
            plt.subplot(4, 3, c+1)
            plt.scatter(x, y, color = col, s = 3)

            plt.tight_layout()
            plt.savefig(PART5_all_classes.png')
            plt.show()
            plt.style.use('default')
```


```
In [9]: # A helper function that trains and returns an MLP classifier with the given hyperparameters.
def Trainer ( arch , lr , train_data , batch_size = 16 , momentum = 0.8 , early_stopping_tol = 4 ) :
    layers = [
        'A' : [ ],
        'B' : [2],
        'C' : [6],
        'D' : [2, 3],
        'E' : [3, 2]
    ]
    input_dim = len(train_data[0][0])
    model_name = 'ARCHITECTURE[{:arch}] + LR {:lr} + str(lr)
    model = Handwritten_Digits_Classifier(model_name, input_dim, layers[arch], 10, lr = lr,
                                           momentum = momentum, early_stopping_tol = early_stopping_tol)
    model_train(train_data, total_epochs = 400, batch_size = batch_size)
    return model
```

```
In [10]: data = Load_Data('optdigits.tra', 'optdigits.tes')
train_data = Split(data, 0.75)

STD. DEVIATION IN THE ATTRIBUTE VALUES
```



DIVERSITY IN THE ATTRIBUTE VALUES



ENTROPY OF THE ATTRIBUTES

PART 02

The following MLP architectures are trained, on the preprocessed dataset, 5 times with the learning rates as 10-1, 10-2, 10-3, 10-4 and 10-5. The length of the list denotes the number of hidden layers and the list itself consists of the number of nodes in each hidden layer (in the respective order from the input layer to the output layer).

(A) Hidden Layers - []

(B) Hidden Layers - [2]

(C) Hidden Layers - [6]

(D) Hidden Layers - [2, 3]

(E) Hidden Layers - [3, 2]

Train 5 MLP models for each of the above architectures, with each one of the above 5 learning rates (25 models in total). Report the various details related to training like training time and number of epochs. Finally evaluate each of the models post-training on the training set and test set and report their respective accuracies and losses.

```
In [13]: _best_lr_for_mlp_architectures__ = dict()
_train_accuracies__ = dict()
_test_accuracies__ = dict()

In [14]: arch_id = 'A'

for lr in [0.1, 0.01, 0.001, 0.0001, 0.00001] :
    print(' * 50, '\n')
    model = Trainer(arch_id, lr, train_data)
    train_acc, test_acc = model_model_analysis_report(train_data, test_data)
    if not arch_id in __train_accuracies__ : __train_accuracies__[arch_id] = dict()
    if not arch_id in __test_accuracies__ : __test_accuracies__[arch_id] = dict()
    __train_accuracies__[arch_id][lr] = train_acc
    __test_accuracies__[arch_id][lr] = test_acc
    print('---' * 50, '\n')

best_lr = None
highest_test_acc_obs = 0.0
for lr, acc in __test_accuracies__[arch_id].items() :
    if acc > highest_test_acc_obs :
        highest_test_acc_obs = acc
        best_lr = lr

__best_lr_for_mlp_architectures__[arch_id] = best_lr
```

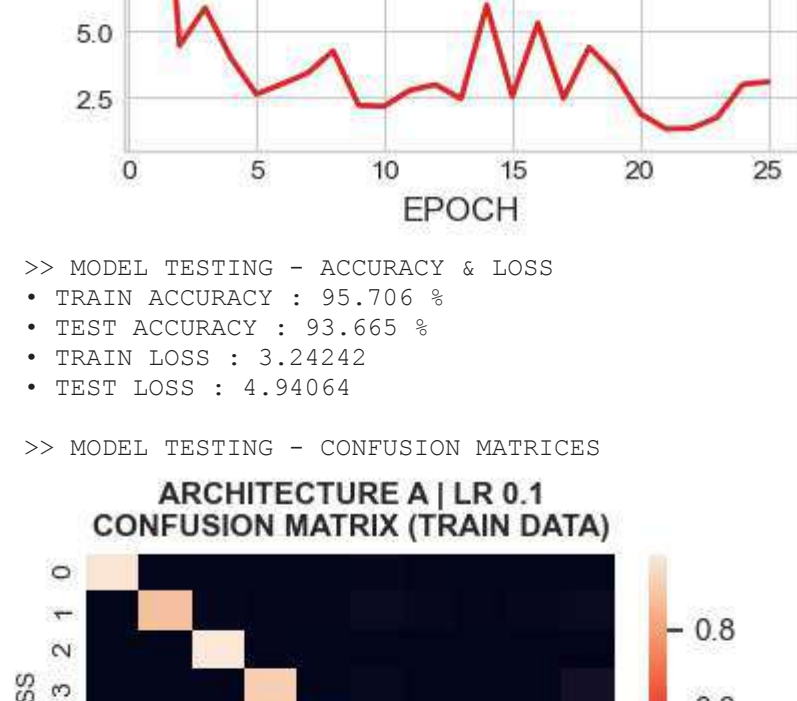
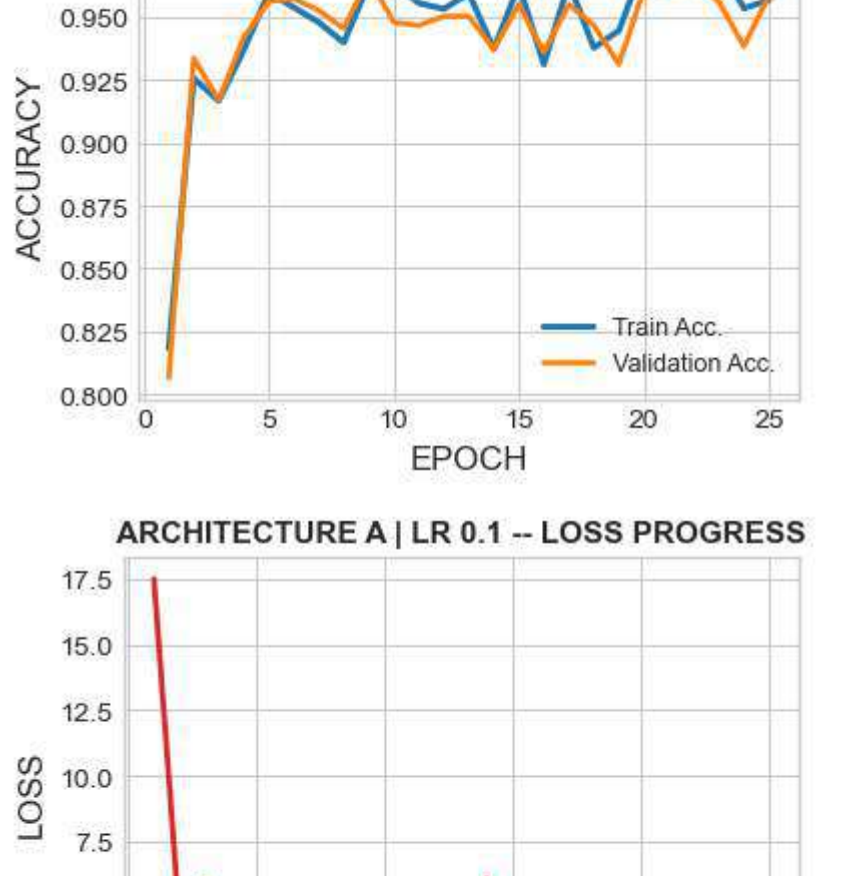

+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.1
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.1
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

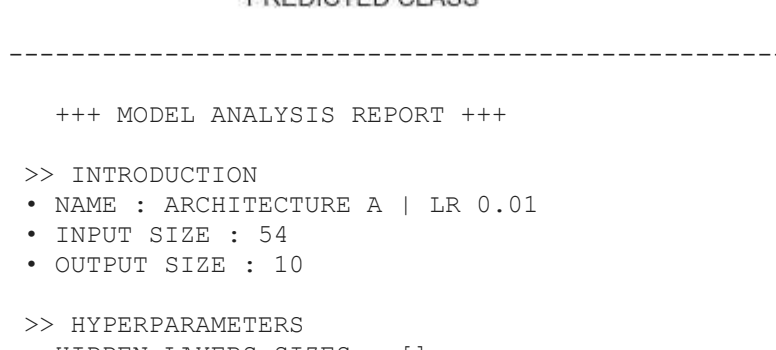
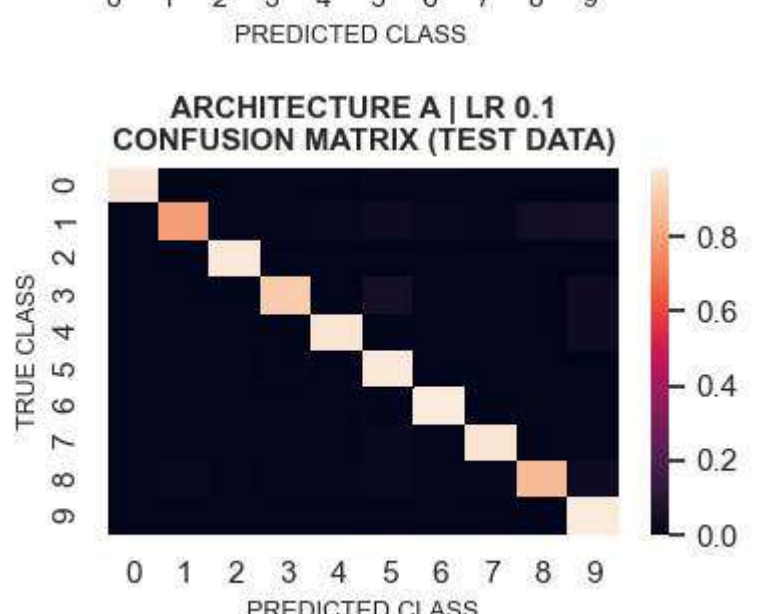
>> TRAINING INFORMATION
• EPOCHS : 25
• TRAINING TIME : 0.131 mins
• STOPPING REASON : [Early Stopping] Training loss was not decreasing
• STOPPING TRAIN LOSS : 3.07463
• STOPPING TRAIN ACCURACY : 95.670 %
• STOPPING VALIDATION ACCURACY : 95.848 %

>> TRAINING PROGRESS



>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 95.706 %
• TEST ACCURACY : 93.665 %
• TRAIN LOSS : 3.24242
• TEST LOSS : 4.94064

>> MODEL TESTING - CONFUSION MATRICES



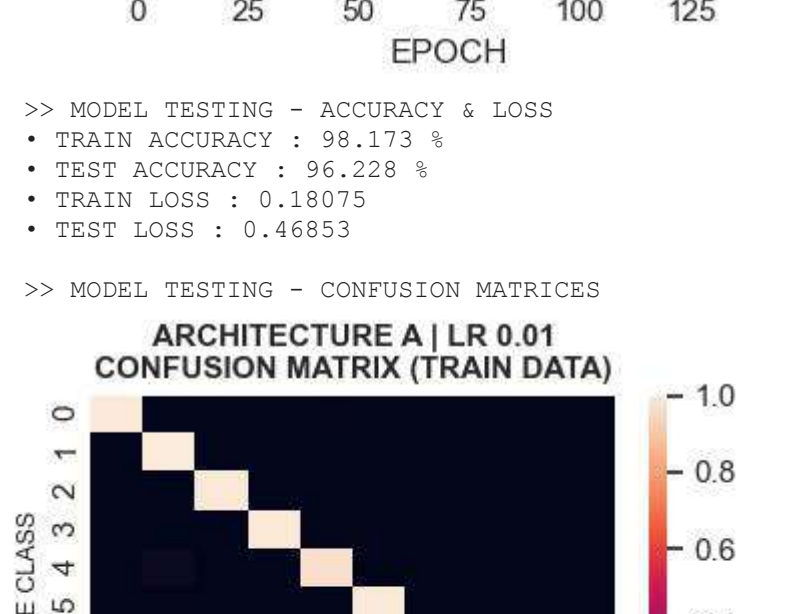

+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.01
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.01
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

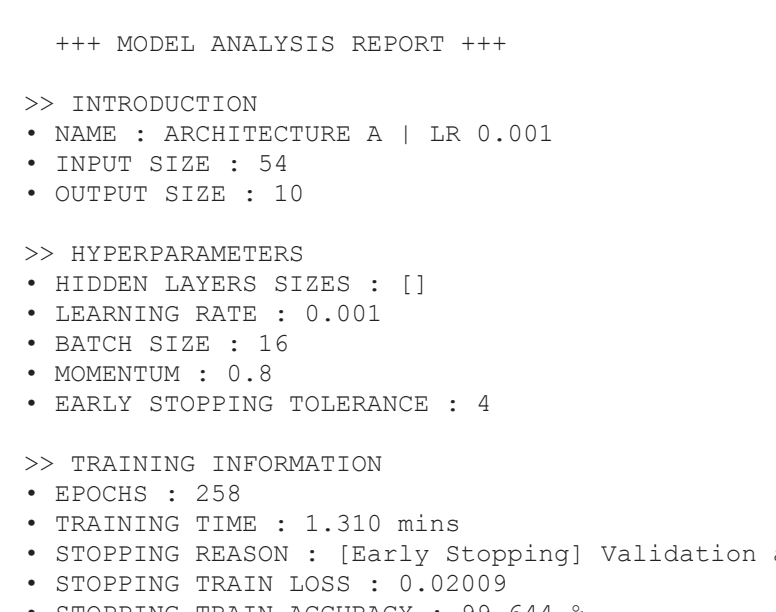
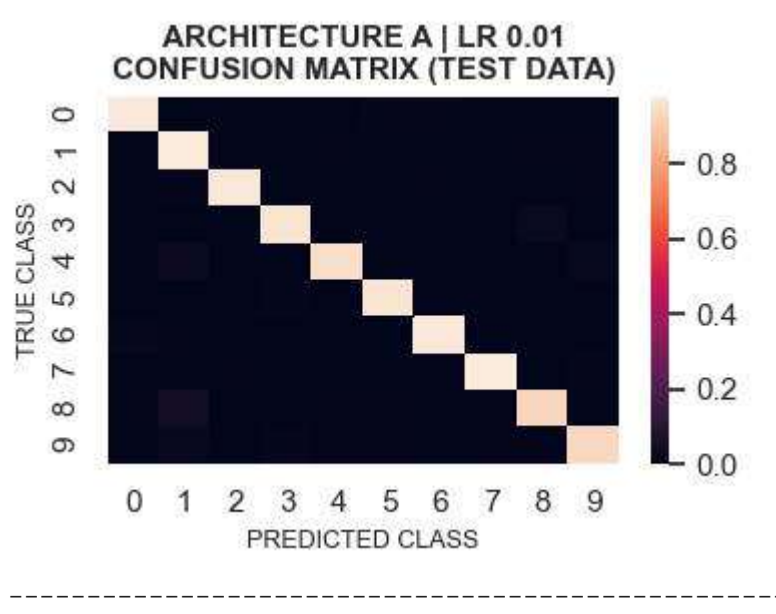
>> TRAINING INFORMATION
• EPOCHS : 140
• TRAINING TIME : 0.715 mins
• STOPPING REASON : [Early Stopping] Training loss was not decreasing
• STOPPING TRAIN LOSS : 0.05690
• STOPPING TRAIN ACCURACY : 98.754 %
• STOPPING VALIDATION ACCURACY : 95.848 %

>> TRAINING PROGRESS



>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 98.173 %
• TEST ACCURACY : 96.228 %
• TRAIN LOSS : 0.18075
• TEST LOSS : 0.46853

>> MODEL TESTING - CONFUSION MATRICES



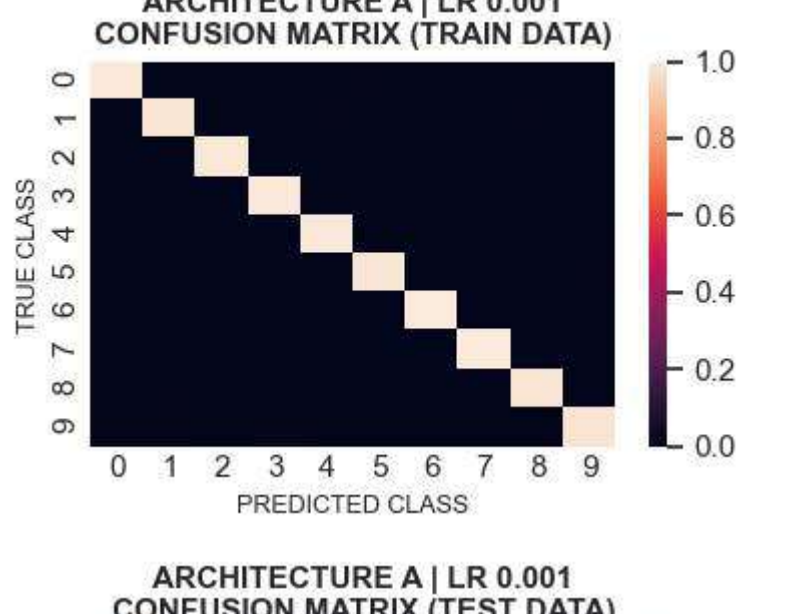

+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

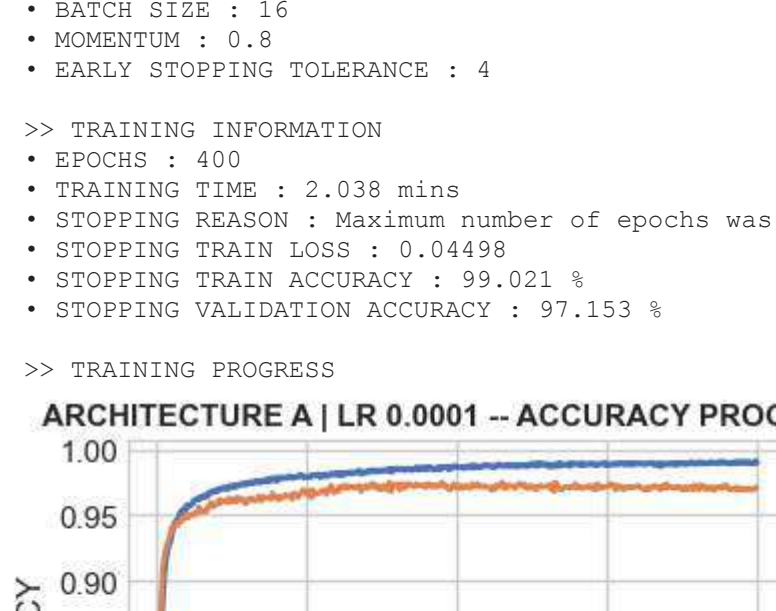
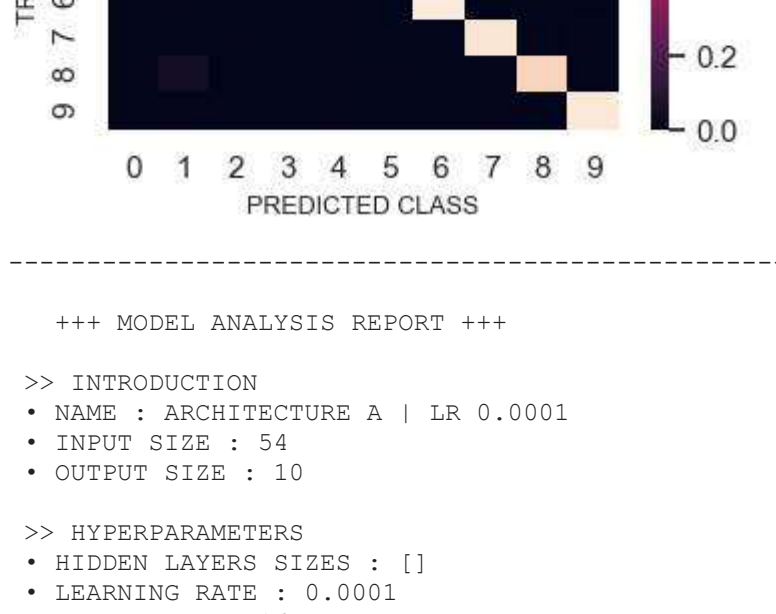
>> TRAINING INFORMATION
• EPOCHS : 258
• TRAINING TIME : 1.310 mins
• STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
• STOPPING TRAIN LOSS : 0.02009
• STOPPING TRAIN ACCURACY : 99.644 %
• STOPPING VALIDATION ACCURACY : 96.679 %

>> TRAINING PROGRESS



>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 99.051 %
• TEST ACCURACY : 96.940 %
• TRAIN LOSS : 0.04211
• TEST LOSS : 0.12488

>> MODEL TESTING - CONFUSION MATRICES



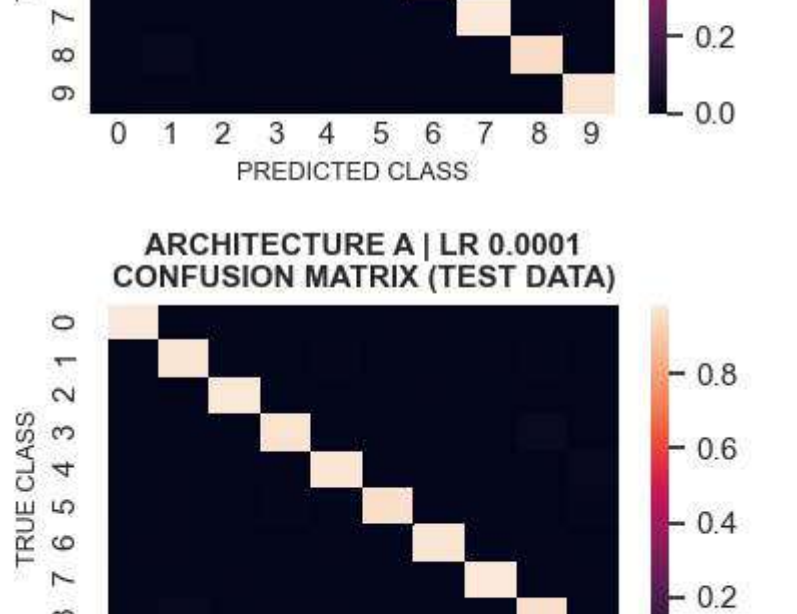
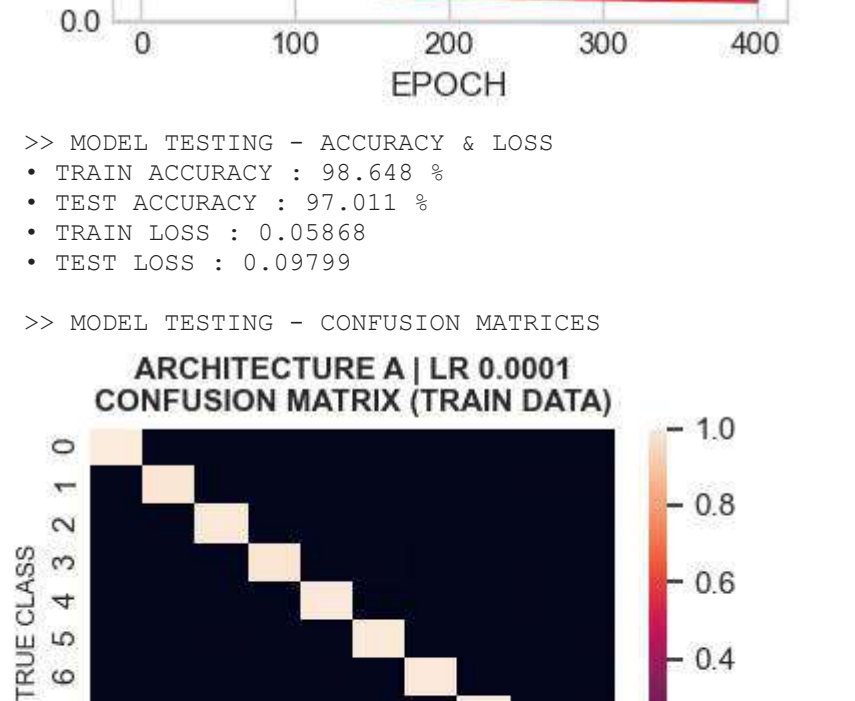
+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4


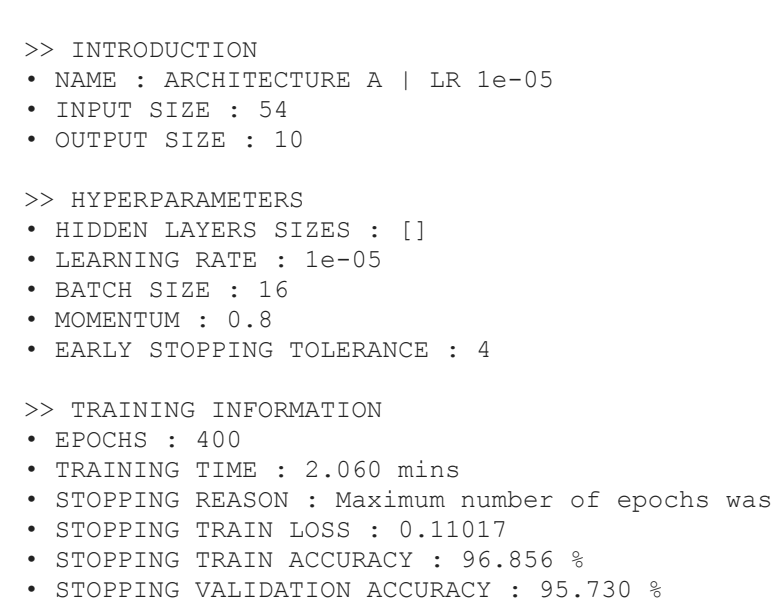
>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 2.038 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 0.04498
• STOPPING TRAIN ACCURACY : 99.021 %
• STOPPING VALIDATION ACCURACY : 97.153 %

>> TRAINING PROGRESS



>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 98.648 %
• TEST ACCURACY : 97.011 %
• TRAIN LOSS : 0.05868
• TEST LOSS : 0.09799

>> MODEL TESTING - CONFUSION MATRICES



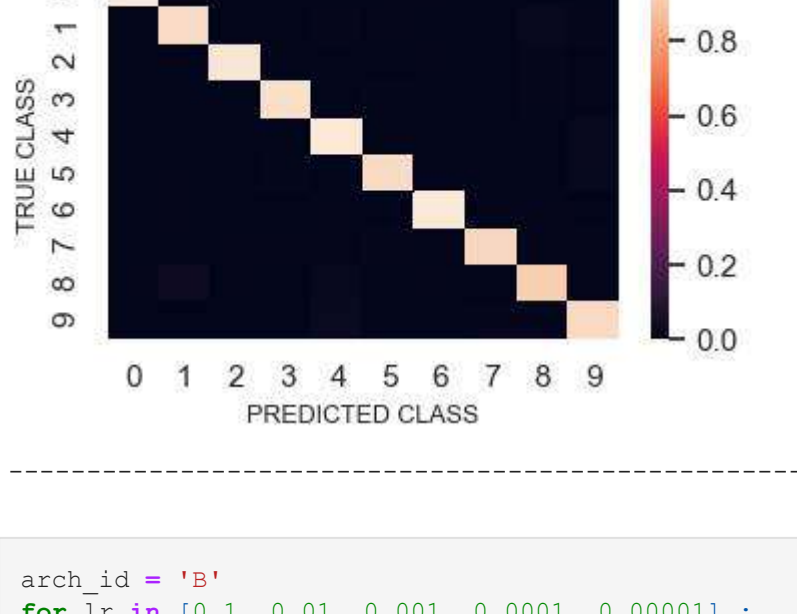

+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 1e-05
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 1e-05
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

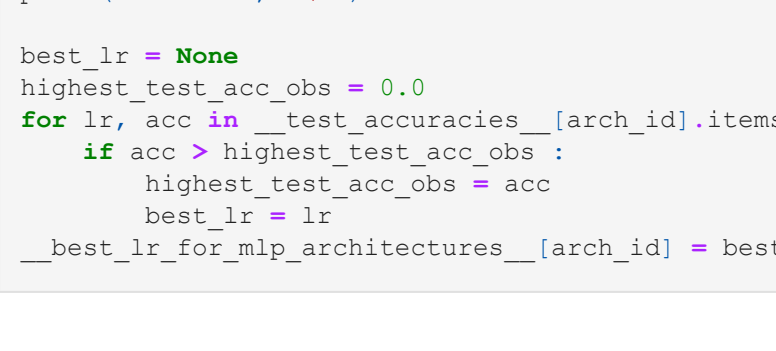
>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 2.060 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 0.11017
• STOPPING TRAIN ACCURACY : 96.856 %
• STOPPING VALIDATION ACCURACY : 95.730 %

>> TRAINING PROGRESS



>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 96.631 %
• TEST ACCURACY : 95.587 %
• TRAIN LOSS : 0.11937
• TEST LOSS : 0.14934

>> MODEL TESTING - CONFUSION MATRICES



```
In [15]: arch_id = 'B'
for lr in [0.1, 0.01, 0.001, 0.0001, 0.00001] :
    print('--' * 50, '\n')
    model = Trainer(arch_id, lr, train_data)
    train_acc, test_acc = model_model_analysis_report(train_data, test_data)
    if not arch_id in train_accuracies : train_accuracies[arch_id] = dict()
    if not arch_id in test_accuracies : test_accuracies[arch_id] = dict()
    train_accuracies[arch_id][lr] = train_acc
    test_accuracies[arch_id][lr] = test_acc
    print('--' * 50, '\n')

best_lr = None
highest_test_acc_obs = 0.0
for lr, acc in test_accuracies[arch_id].items() :
    if acc > highest_test_acc_obs :
        highest_test_acc_obs = acc
        best_lr = lr
_best_lr_for_mlp_architectures__[arch_id] = best_lr
```


+++ MODEL ANALYSIS REPORT +++

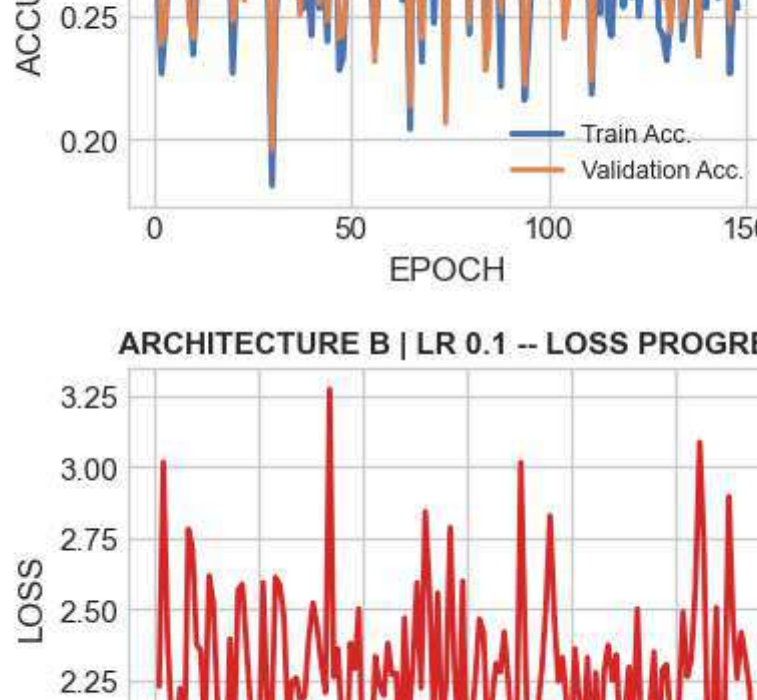
```
>> INTRODUCTION
• NAME : ARCHITECTURE B | LR 0.1
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : [2]
• LEARNING RATE : 0.1
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

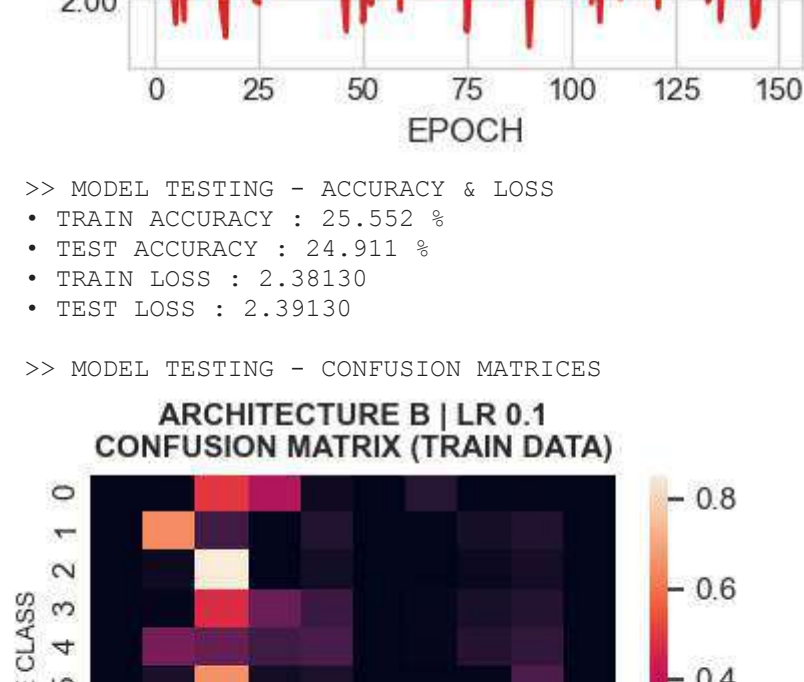
>> TRAINING INFORMATION
• EPOCHS : 145
• TRAINING TIME : 1.006 mins
• STOPPING REASON : [Early Stopping] Training loss was not decreasing
• STOPPING TRAIN LOSS : 2.40891
• STOPPING TRAIN ACCURACY : 25.326 %
• STOPPING VALIDATION ACCURACY : 26.453 %

>> TRAINING PROGRESS
```

ARCHITECTURE B | LR 0.1 -- ACCURACY PROGRESS




ARCHITECTURE B | LR 0.1 -- LOSS PROGRESS



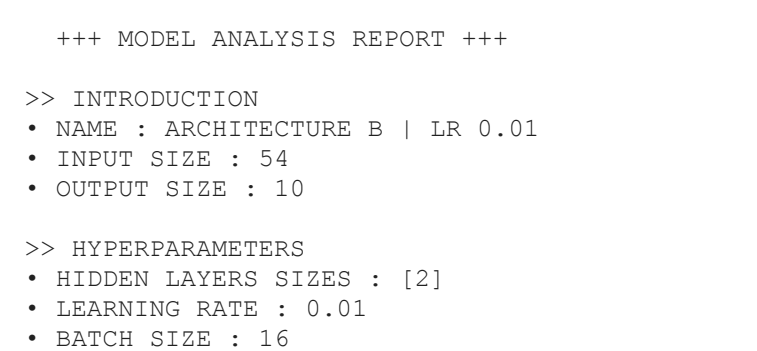
```
>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 25.552 %
• TEST ACCURACY : 24.911 %
• TRAIN LOSS : 2.38130
• TEST LOSS : 2.39130
```

```
>> MODEL TESTING - CONFUSION MATRICES
```

ARCHITECTURE B | LR 0.1
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE B | LR 0.1
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++


```
>> INTRODUCTION
• NAME : ARCHITECTURE B | LR 0.01
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : [2]
• LEARNING RATE : 0.01
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

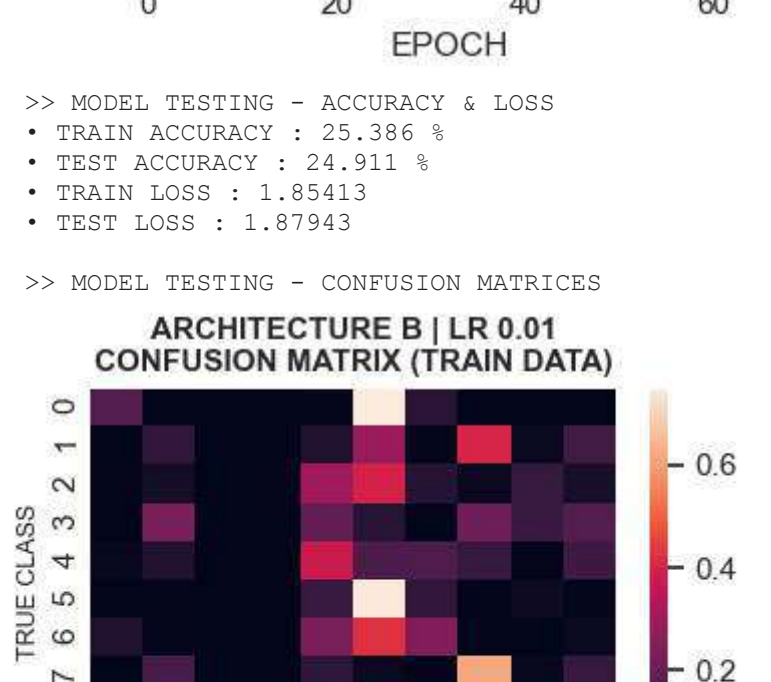
>> TRAINING INFORMATION
• EPOCHS : 63
• TRAINING TIME : 0.417 mins
• STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
• STOPPING TRAIN LOSS : 1.86005
• STOPPING TRAIN ACCURACY : 25.326 %
• STOPPING VALIDATION ACCURACY : 25.623 %

>> TRAINING PROGRESS
```

ARCHITECTURE B | LR 0.01 -- ACCURACY PROGRESS




ARCHITECTURE B | LR 0.01 -- LOSS PROGRESS



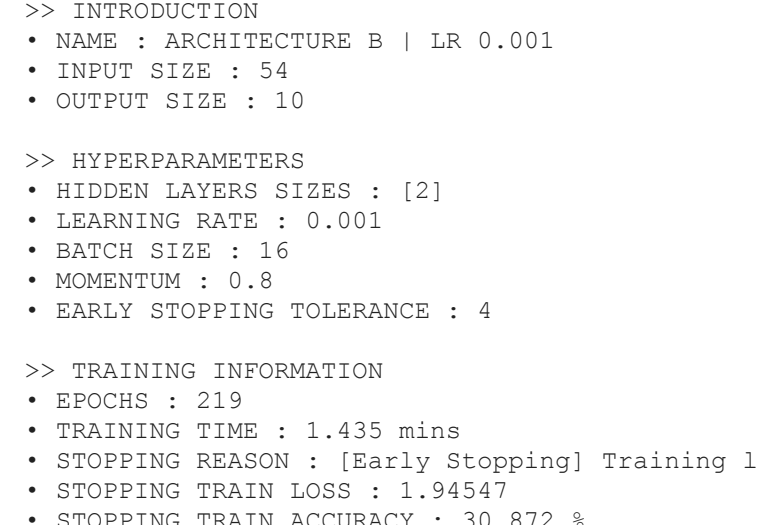
```
>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 25.386 %
• TEST ACCURACY : 24.911 %
• TRAIN LOSS : 1.85413
• TEST LOSS : 1.87943
```

```
>> MODEL TESTING - CONFUSION MATRICES
```

ARCHITECTURE B | LR 0.01
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE B | LR 0.01
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++


```
>> INTRODUCTION
• NAME : ARCHITECTURE B | LR 0.001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : [2]
• LEARNING RATE : 0.001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

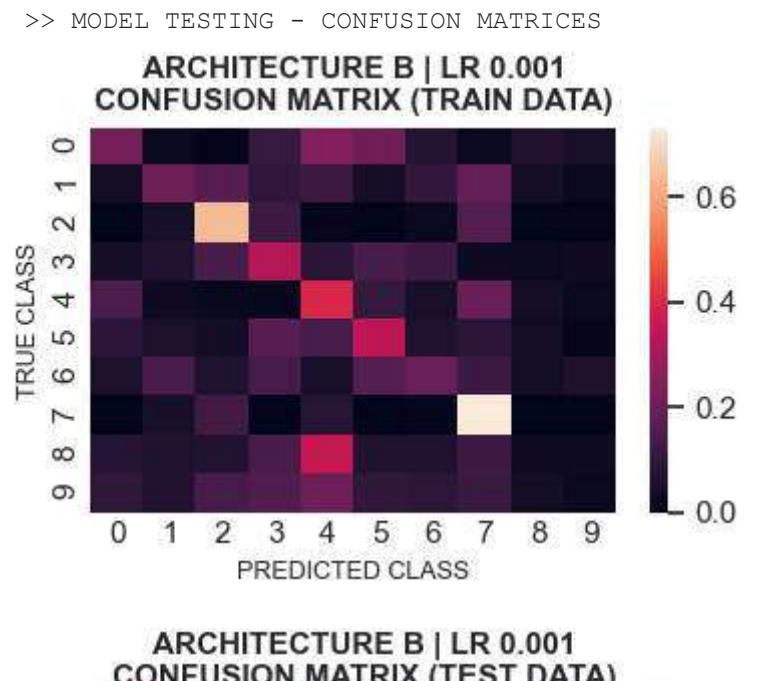
>> TRAINING INFORMATION
• EPOCHS : 219
• TRAINING TIME : 1.435 mins
• STOPPING REASON : [Early Stopping] Training loss was not decreasing
• STOPPING TRAIN LOSS : 1.94547
• STOPPING TRAIN ACCURACY : 30.872 %
• STOPPING VALIDATION ACCURACY : 31.791 %

>> TRAINING PROGRESS
```

ARCHITECTURE B | LR 0.001 -- ACCURACY PROGRESS



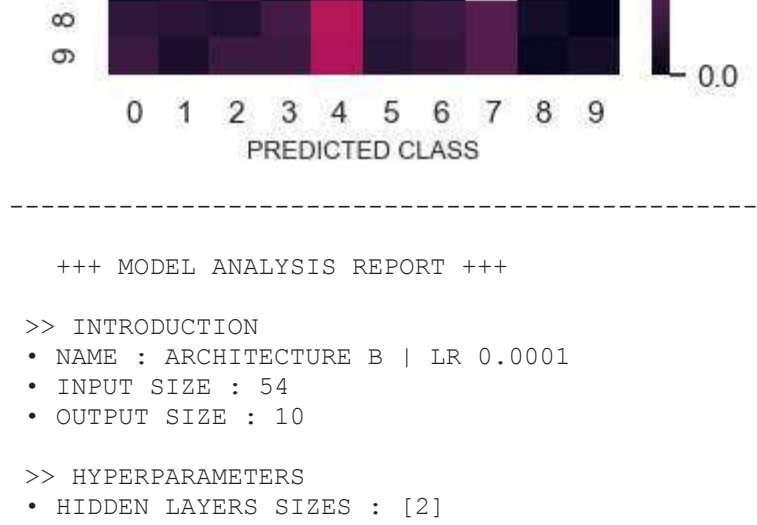
ARCHITECTURE B | LR 0.001 -- LOSS PROGRESS



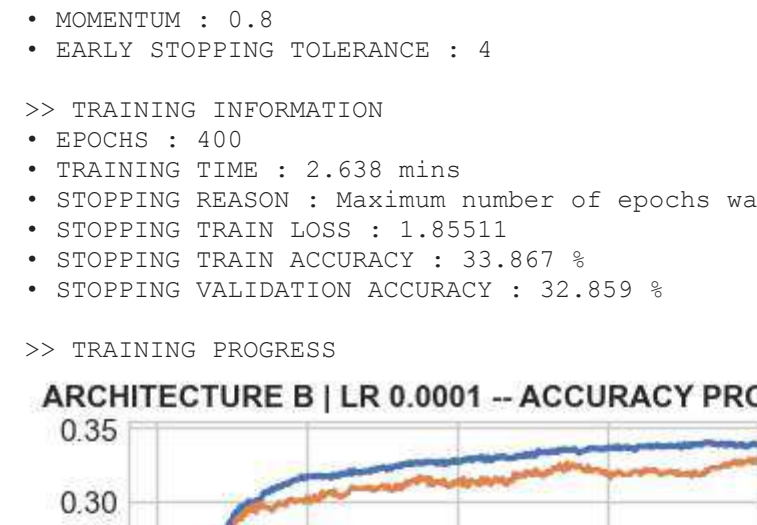
```
>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 31.056 %
• TEST ACCURACY : 30.391 %
• TRAIN LOSS : 1.94052
• TEST LOSS : 1.96083
```

```
>> MODEL TESTING - CONFUSION MATRICES
```

ARCHITECTURE B | LR 0.001
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE B | LR 0.001
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

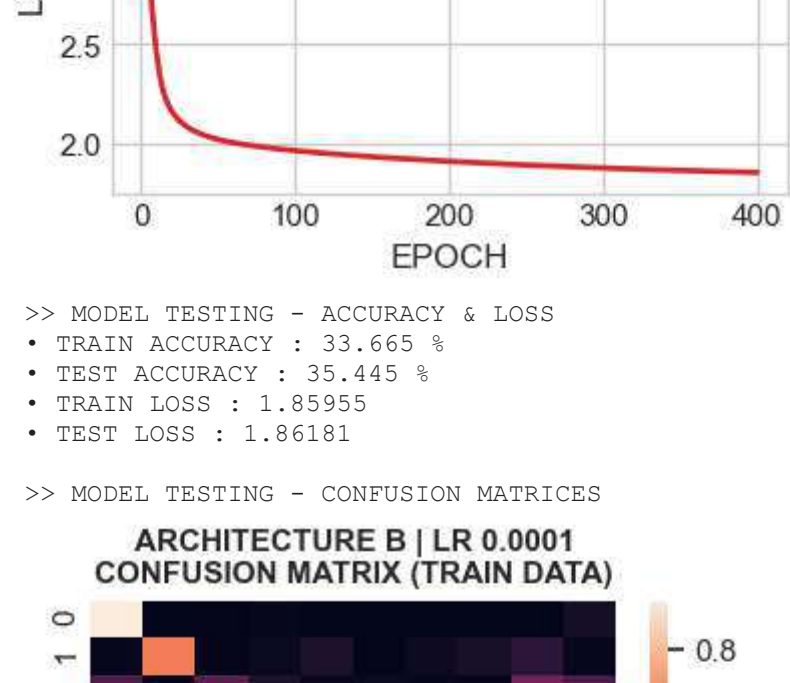
```
>> INTRODUCTION
• NAME : ARCHITECTURE B | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : [2]
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

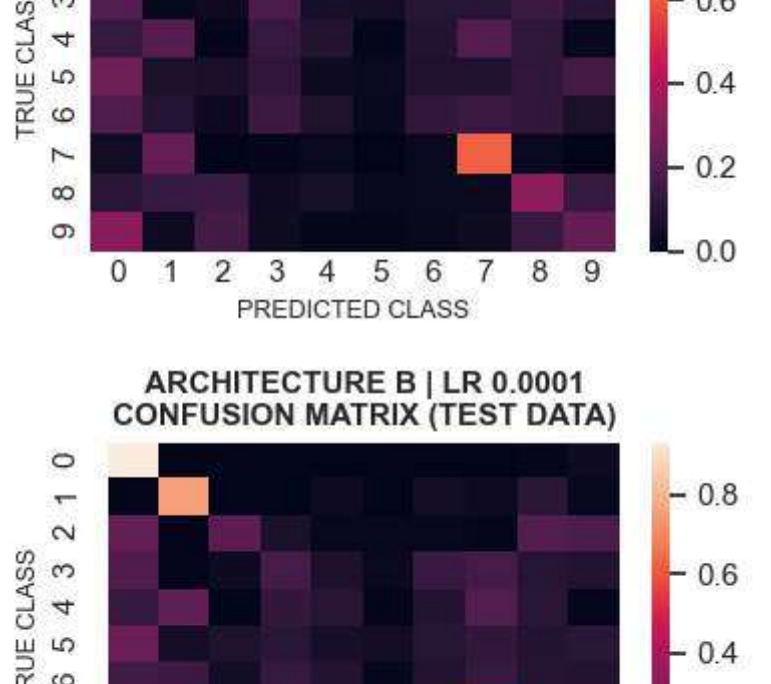
>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 2.638 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 1.85511
• STOPPING TRAIN ACCURACY : 33.867 %
• STOPPING VALIDATION ACCURACY : 32.859 %

>> TRAINING PROGRESS
```

ARCHITECTURE B | LR 0.0001 -- ACCURACY PROGRESS



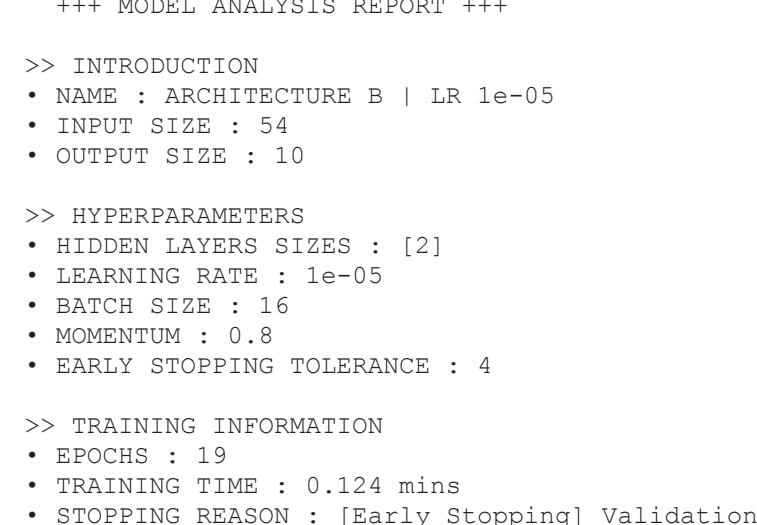
ARCHITECTURE B | LR 0.0001 -- LOSS PROGRESS




```
>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 33.665 %
• TEST ACCURACY : 35.445 %
• TRAIN LOSS : 1.85955
• TEST LOSS : 1.86181
```

```
>> MODEL TESTING - CONFUSION MATRICES
```

ARCHITECTURE B | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE B | LR 0.0001
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

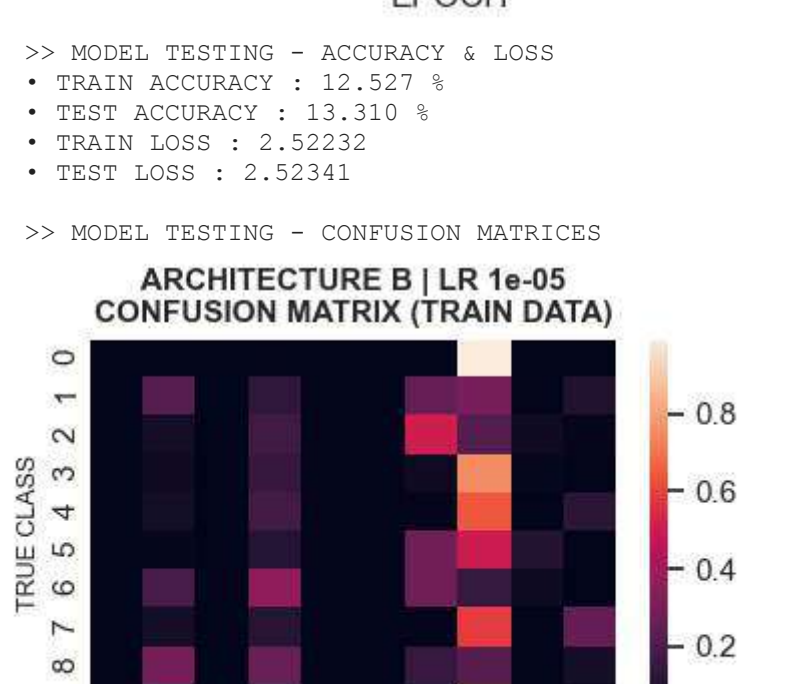
```
>> INTRODUCTION
• NAME : ARCHITECTURE B | LR 1e-05
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : [2]
• LEARNING RATE : 1e-05
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4


>> TRAINING INFORMATION
• EPOCHS : 19
• TRAINING TIME : 0.124 mins
• STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
• STOPPING TRAIN LOSS : 2.52213
• STOPPING TRAIN ACCURACY : 12.574 %
• STOPPING VALIDATION ACCURACY : 12.337 %

>> TRAINING PROGRESS
```

ARCHITECTURE B | LR 1e-05 -- ACCURACY PROGRESS



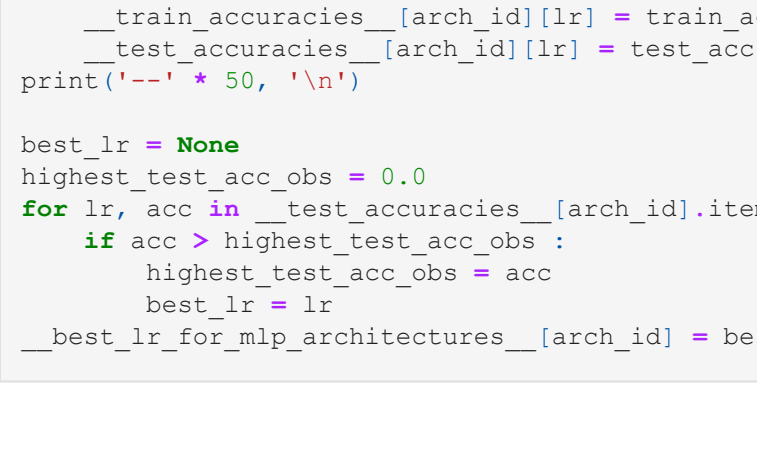
ARCHITECTURE B | LR 1e-05 -- LOSS PROGRESS



```
>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 12.527 %
• TEST ACCURACY : 13.310 %
• TRAIN LOSS : 2.52232
• TEST LOSS : 2.52341
```

```
>> MODEL TESTING - CONFUSION MATRICES
```

ARCHITECTURE B | LR 1e-05
CONFUSION MATRIX (TRAIN DATA)




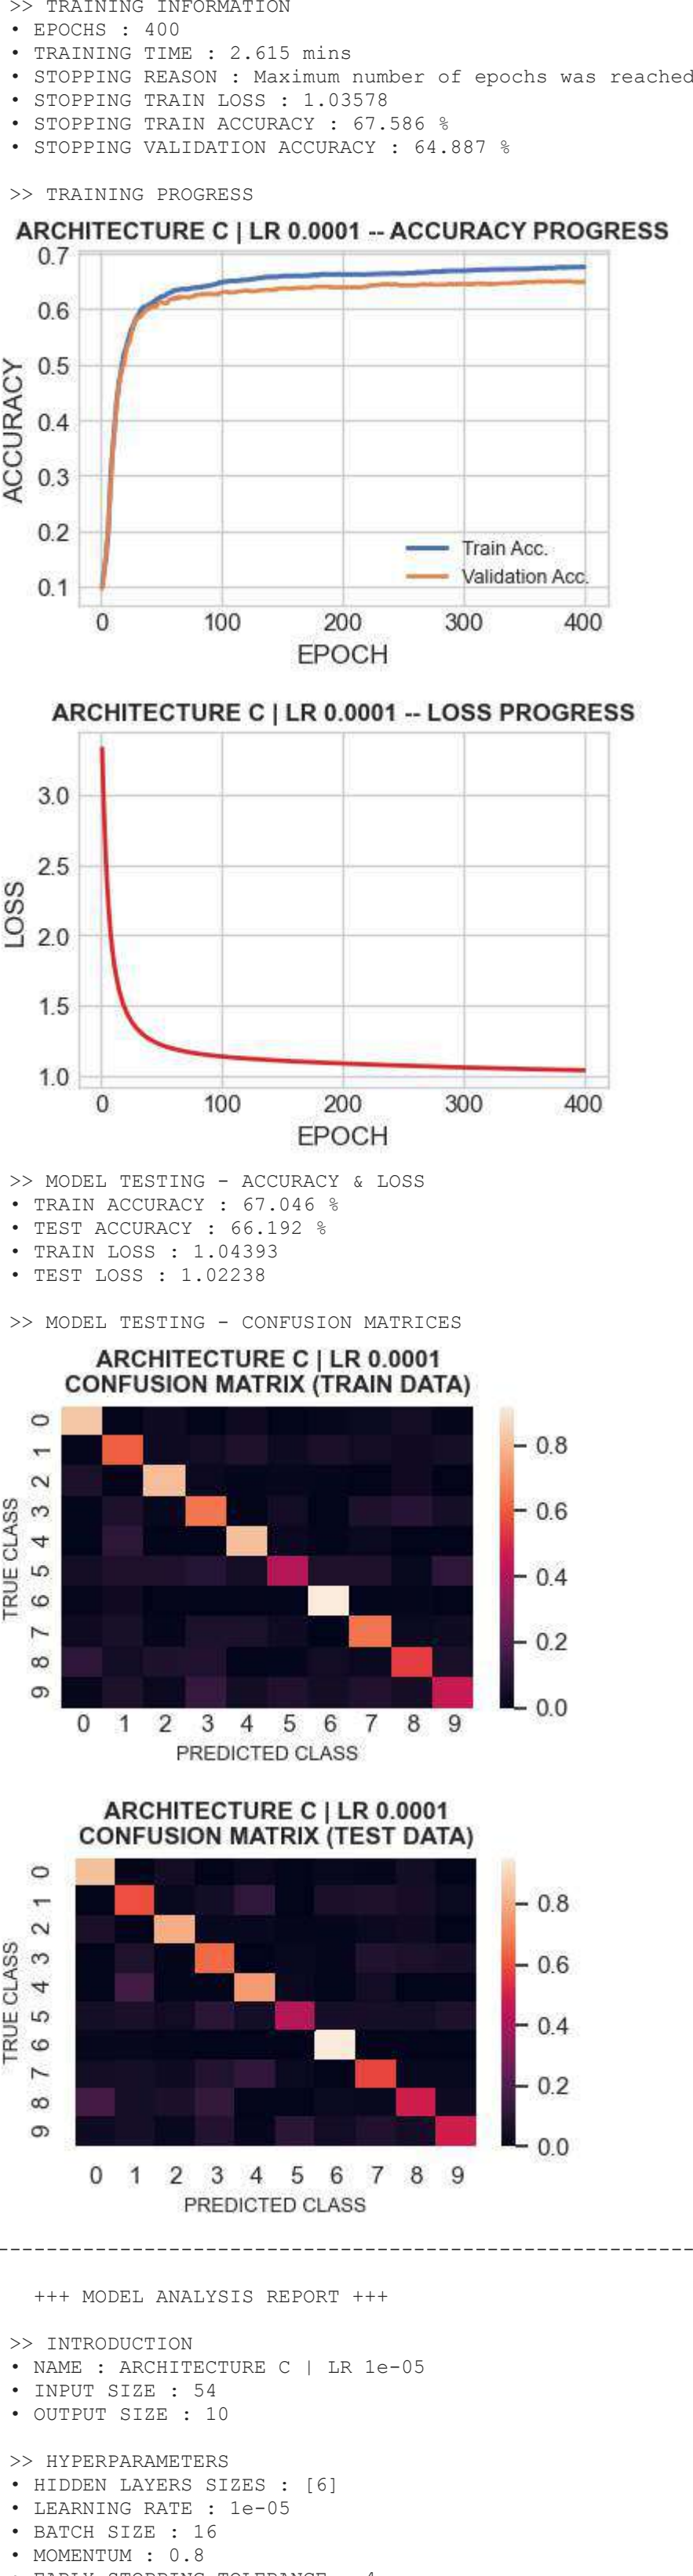
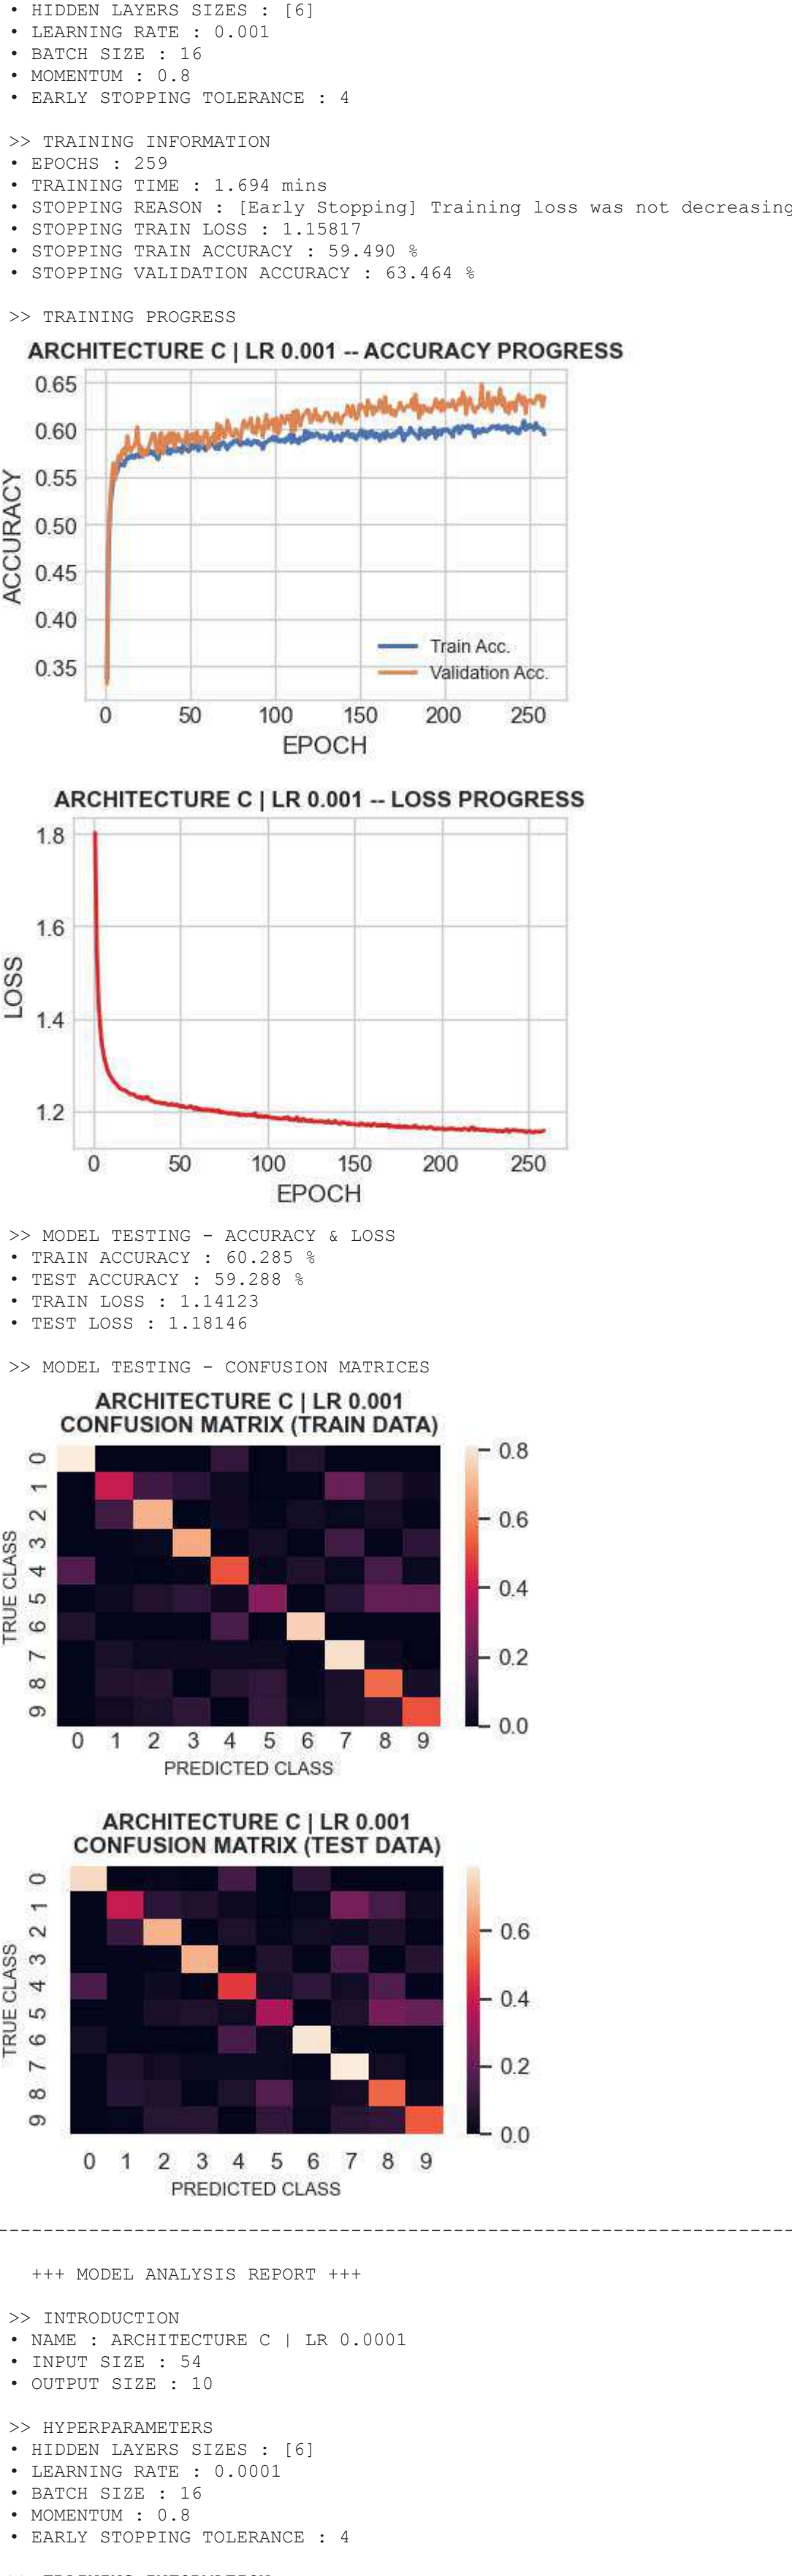
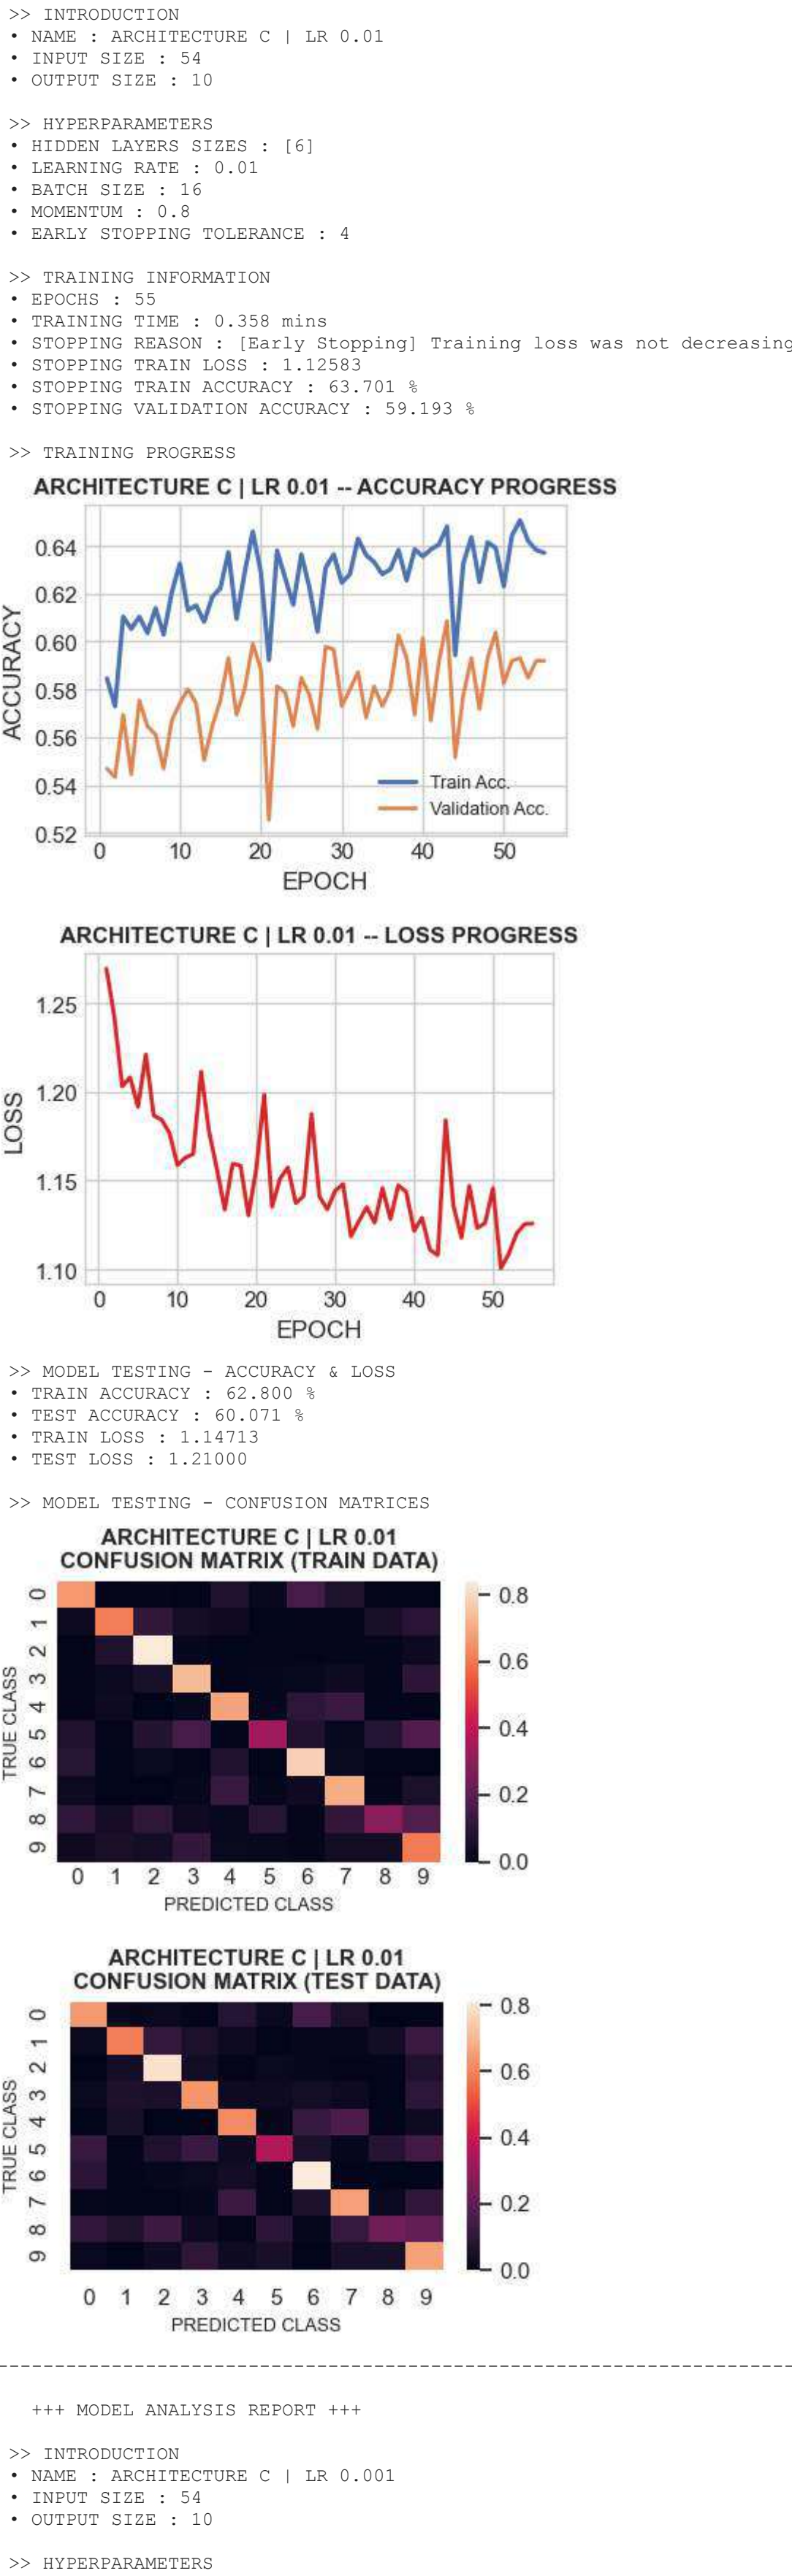
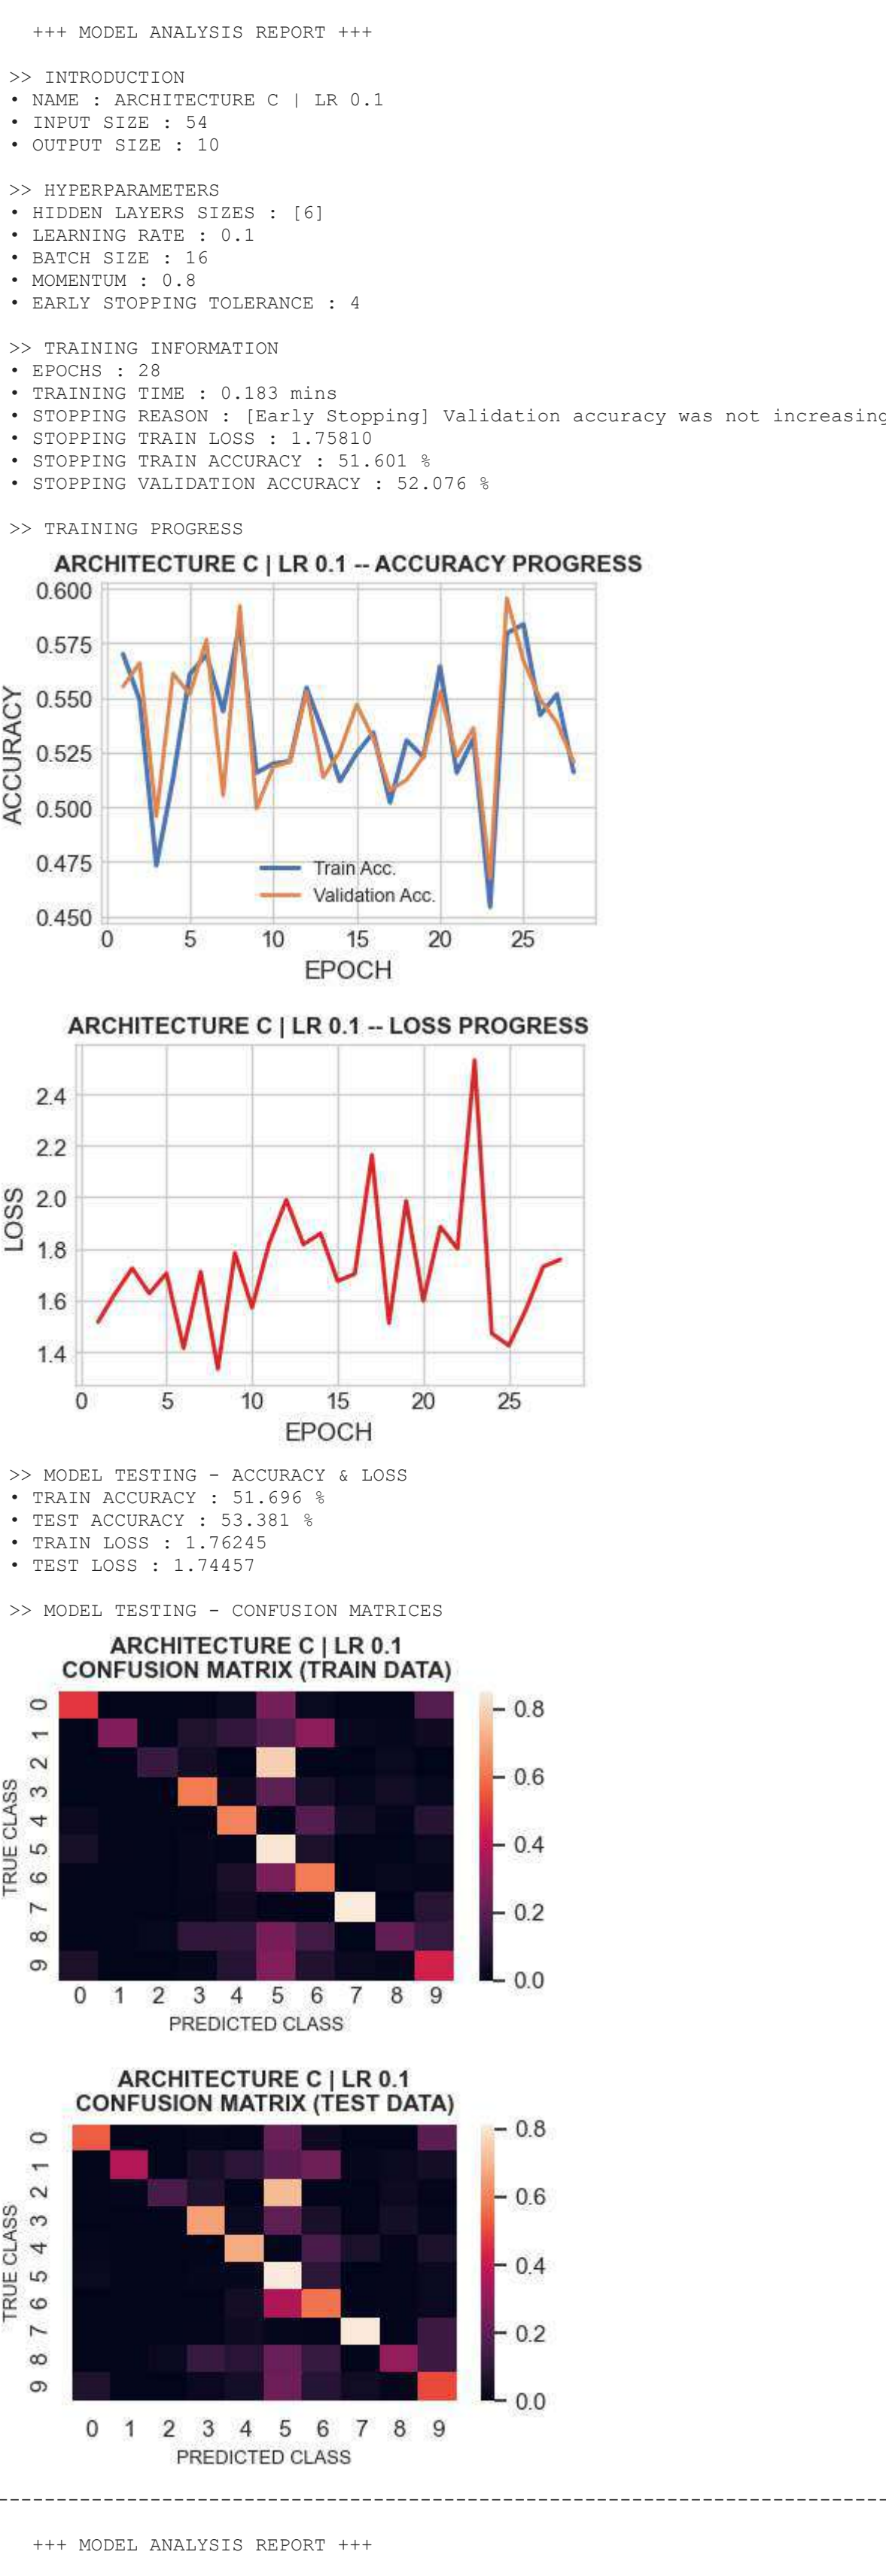
ARCHITECTURE B | LR 1e-05
CONFUSION MATRIX (TEST DATA)

In [16]:

```
arch_id = 'c'
for lr in [0.1, 0.01, 0.001, 0.0001, 0.00001]:
    print('--> * %s' % lr)
    model = Trainer(arch_id, lr, train_data)
    train_acc, test_acc = model.model_analysis_report(train_data, test_data)
    if not arch_id in train_accuracies_: train_accuracies_[arch_id] = dict()
    if not arch_id in test_accuracies_: test_accuracies_[arch_id] = dict()
    train_accuracies_[arch_id][lr] = train_acc
    test_accuracies_[arch_id][lr] = test_acc
    print('--> * %s' % lr)

best_lr = None
highest_test_acc_obs = 0.0
for lr, acc in test_accuracies_[arch_id].items():
    if acc > highest_test_acc_obs:
        highest_test_acc_obs = acc
        best_lr = lr

_best_lr_for_mlp_architectures_[arch_id] = best_lr
```

In [17]:

```
arch_id = 'D'
for lr in [0.1, 0.01, 0.001, 0.0001, 0.00001] :
    print('--' * 50, '\n')
    model = Trainer(arch_id, lr, train_data)
    train_acc, test_acc = model.model_analysis_report(train_data, test_data)
    if not arch_id in _train_accuracies_ : _train_accuracies_[arch_id] = dict()
    if not arch_id in _test_accuracies_ : _test_accuracies_[arch_id] = dict()
    _train_accuracies_[arch_id][lr] = train_acc
    _test_accuracies_[arch_id][lr] = test_acc
    print('--' * 50, '\n')

best_lr = None
highest_test_acc_obs = 0.0
for lr, acc in _test_accuracies_.items() :
    if acc > highest_test_acc_obs :
        highest_test_acc_obs = acc
        best_lr = lr
_best_lr_for_mlp_architectures_[arch_id] = best_lr
```


+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 0.1
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS


- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.1
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION

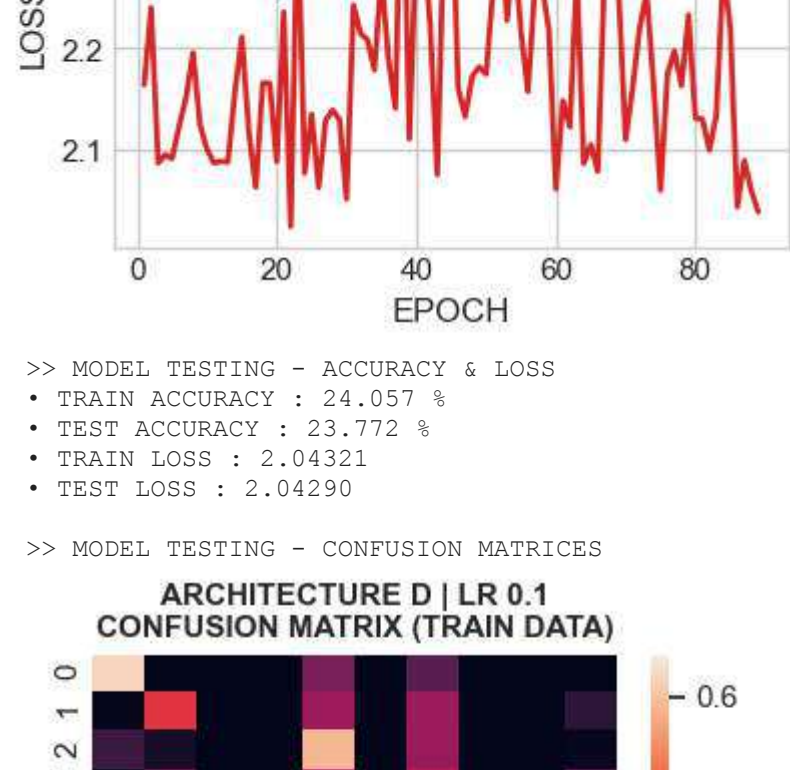
- EPOCHS : 89
- TRAINING TIME : 0.712 mins
- STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
- STOPPING TRAIN LOSS : 2.03772
- STOPPING TRAIN ACCURACY : 24.614 %
- STOPPING VALIDATION ACCURACY : 21.827 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.1 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 0.1 -- LOSS PROGRESS




>> MODEL TESTING - ACCURACY & LOSS

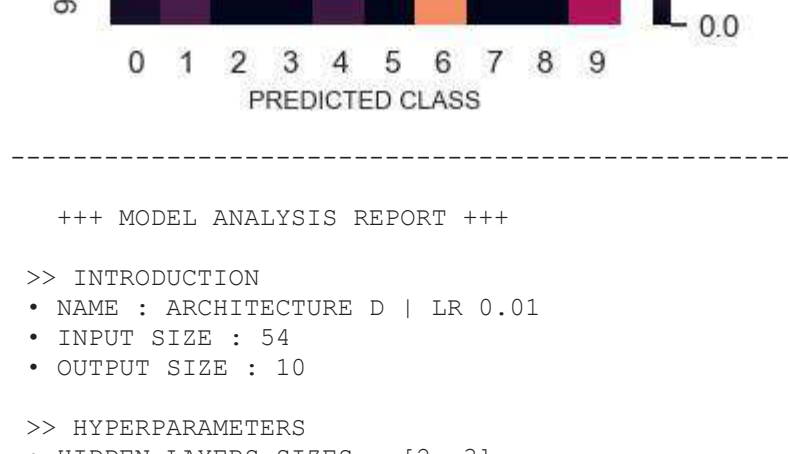
- TRAIN ACCURACY : 24.057 %
- TEST ACCURACY : 23.772 %
- TRAIN LOSS : 2.04321
- TEST LOSS : 2.04290

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.1
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE D | LR 0.1
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 0.01
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS


- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.01
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION

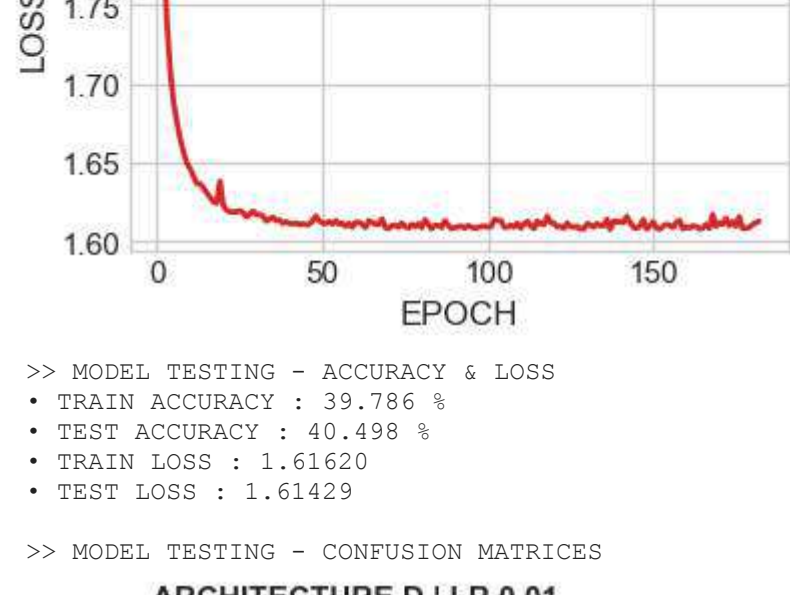
- EPOCHS : 182
- TRAINING TIME : 1.474 mins
- STOPPING REASON : [Early Stopping] Training loss was not decreasing
- STOPPING TRAIN LOSS : 1.61253
- STOPPING TRAIN ACCURACY : 40.095 %
- STOPPING VALIDATION ACCURACY : 38.553 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.01 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 0.01 -- LOSS PROGRESS

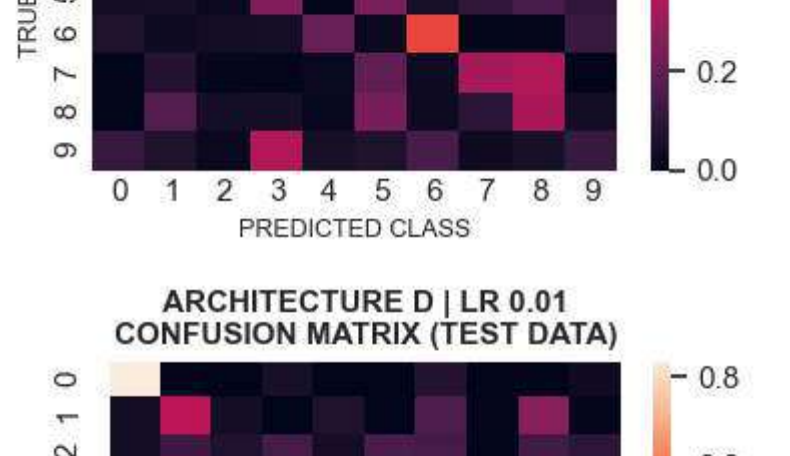


>> MODEL TESTING - ACCURACY & LOSS


- TRAIN ACCURACY : 39.786 %
- TEST ACCURACY : 40.498 %
- TRAIN LOSS : 1.61620
- TEST LOSS : 1.61429

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.01
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE D | LR 0.01
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 0.001
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS


- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.001
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION


- EPOCHS : 195
- TRAINING TIME : 1.576 mins
- STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
- STOPPING TRAIN LOSS : 1.88739
- STOPPING TRAIN ACCURACY : 31.168 %
- STOPPING VALIDATION ACCURACY : 32.384 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.001 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 0.001 -- LOSS PROGRESS




>> MODEL TESTING - ACCURACY & LOSS


- TRAIN ACCURACY : 31.412 %
- TEST ACCURACY : 32.242 %
- TRAIN LOSS : 1.89160
- TEST LOSS : 1.90526

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.001
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE D | LR 0.001
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 0.0001
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS


- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.0001
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION


- EPOCHS : 400
- TRAINING TIME : 3.226 mins
- STOPPING REASON : Maximum number of epochs was reached
- STOPPING TRAIN LOSS : 1.95640
- STOPPING TRAIN ACCURACY : 25.860 %
- STOPPING VALIDATION ACCURACY : 22.183 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.0001 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 0.0001 -- LOSS PROGRESS



>> MODEL TESTING - ACCURACY & LOSS

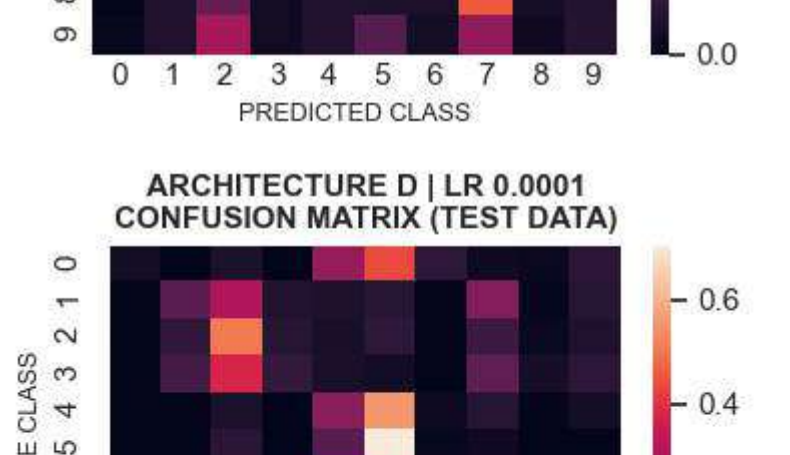
- TRAIN ACCURACY : 25.125 %
- TEST ACCURACY : 24.128 %
- TRAIN LOSS : 1.96389
- TEST LOSS : 1.97660

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 1e-05
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS

- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 1e-05
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION


- EPOCHS : 36
- TRAINING TIME : 0.286 mins
- STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
- STOPPING TRAIN LOSS : 2.39792
- STOPPING TRAIN ACCURACY : 16.815 %
- STOPPING VALIDATION ACCURACY : 16.251 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 1e-05 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 1e-05 -- LOSS PROGRESS

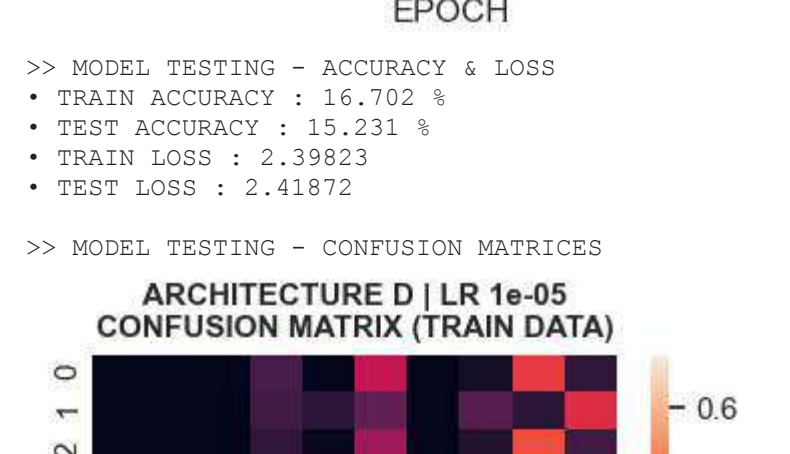


>> MODEL TESTING - ACCURACY & LOSS


- TRAIN ACCURACY : 16.702 %
- TEST ACCURACY : 15.231 %
- TRAIN LOSS : 2.39823
- TEST LOSS : 2.41872

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 1e-05
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE D | LR 1e-05
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 0.0001
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS

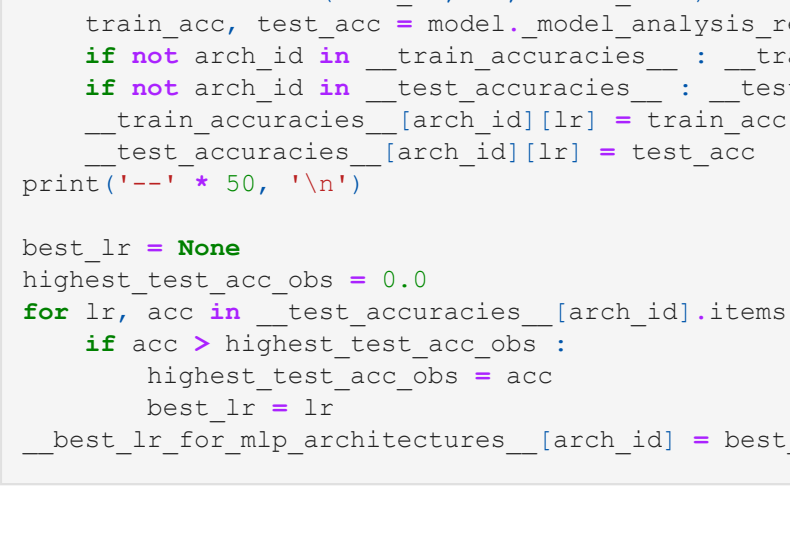
- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.0001
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION

- EPOCHS : 400
- TRAINING TIME : 3.226 mins
- STOPPING REASON : Maximum number of epochs was reached
- STOPPING TRAIN LOSS : 1.95640
- STOPPING TRAIN ACCURACY : 25.860 %
- STOPPING VALIDATION ACCURACY : 22.183 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.0001 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 0.0001 -- LOSS PROGRESS

>> MODEL TESTING - ACCURACY & LOSS

- TRAIN ACCURACY : 25.125 %
- TEST ACCURACY : 24.128 %
- TRAIN LOSS : 1.96389
- TEST LOSS : 1.97660

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)

ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TEST DATA)

+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 0.0001
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS


- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.0001
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION

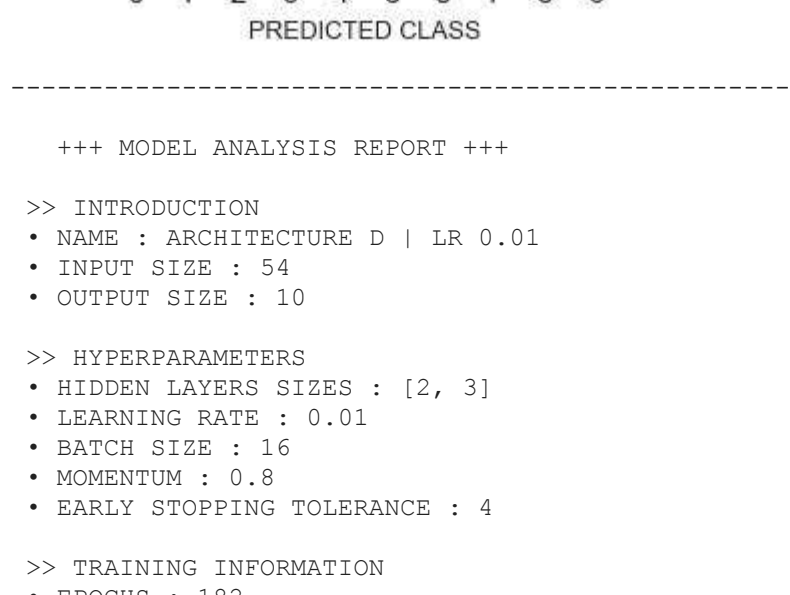
- EPOCHS : 400
- TRAINING TIME : 3.226 mins
- STOPPING REASON : Maximum number of epochs was reached
- STOPPING TRAIN LOSS : 1.95640
- STOPPING TRAIN ACCURACY : 25.860 %
- STOPPING VALIDATION ACCURACY : 22.183 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.0001 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 0.0001 -- LOSS PROGRESS




>> MODEL TESTING - ACCURACY & LOSS


- TRAIN ACCURACY : 25.125 %
- TEST ACCURACY : 24.128 %
- TRAIN LOSS : 1.96389
- TEST LOSS : 1.97660

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

- NAME : ARCHITECTURE D | LR 0.0001
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS


- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.0001
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION

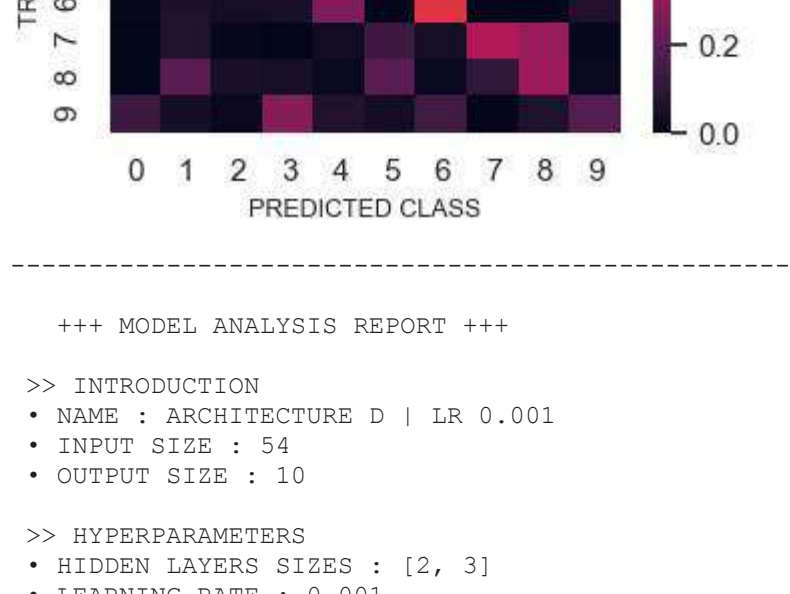
- EPOCHS : 400
- TRAINING TIME : 3.226 mins
- STOPPING REASON : Maximum number of epochs was reached
- STOPPING TRAIN LOSS : 1.95640
- STOPPING TRAIN ACCURACY : 25.860 %
- STOPPING VALIDATION ACCURACY : 22.183 %

>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.0001 -- ACCURACY PROGRESS



ARCHITECTURE D | LR 0.0001 -- LOSS PROGRESS




>> MODEL TESTING - ACCURACY & LOSS


- TRAIN ACCURACY : 25.125 %
- TEST ACCURACY : 24.128 %
- TRAIN LOSS : 1.96389
- TEST LOSS : 1.97660

>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)



ARCHITECTURE D | LR 0.0001
CONFUSION MATRIX (TEST DATA)



+++ MODEL ANALYSIS REPORT +++

>> INTRODUCTION

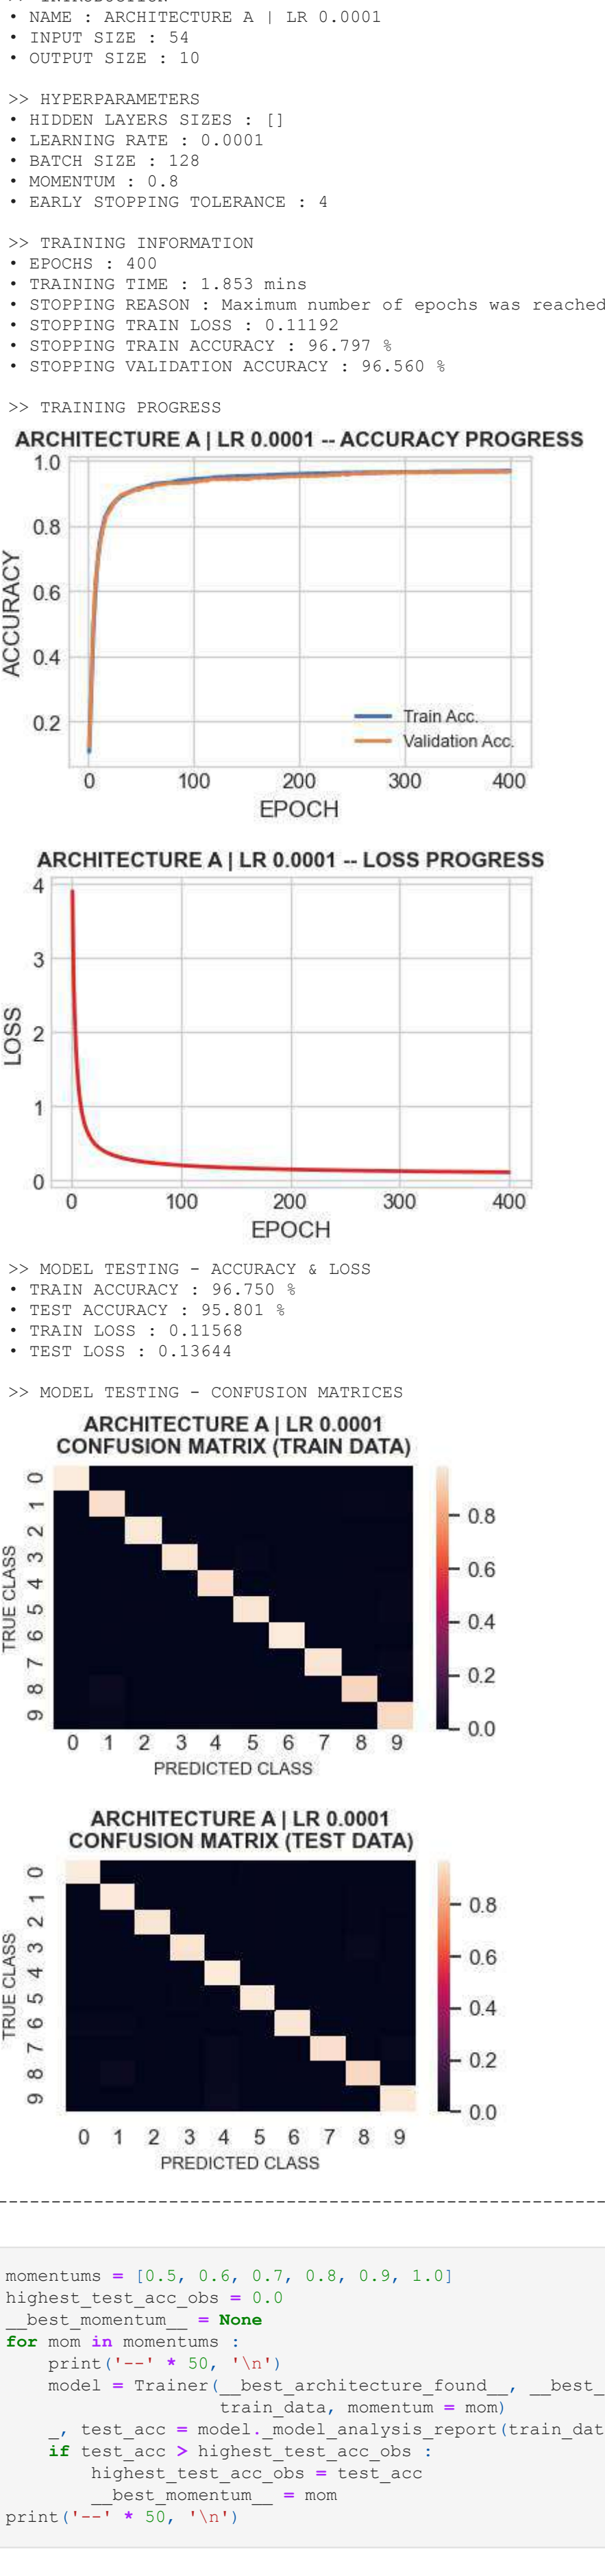
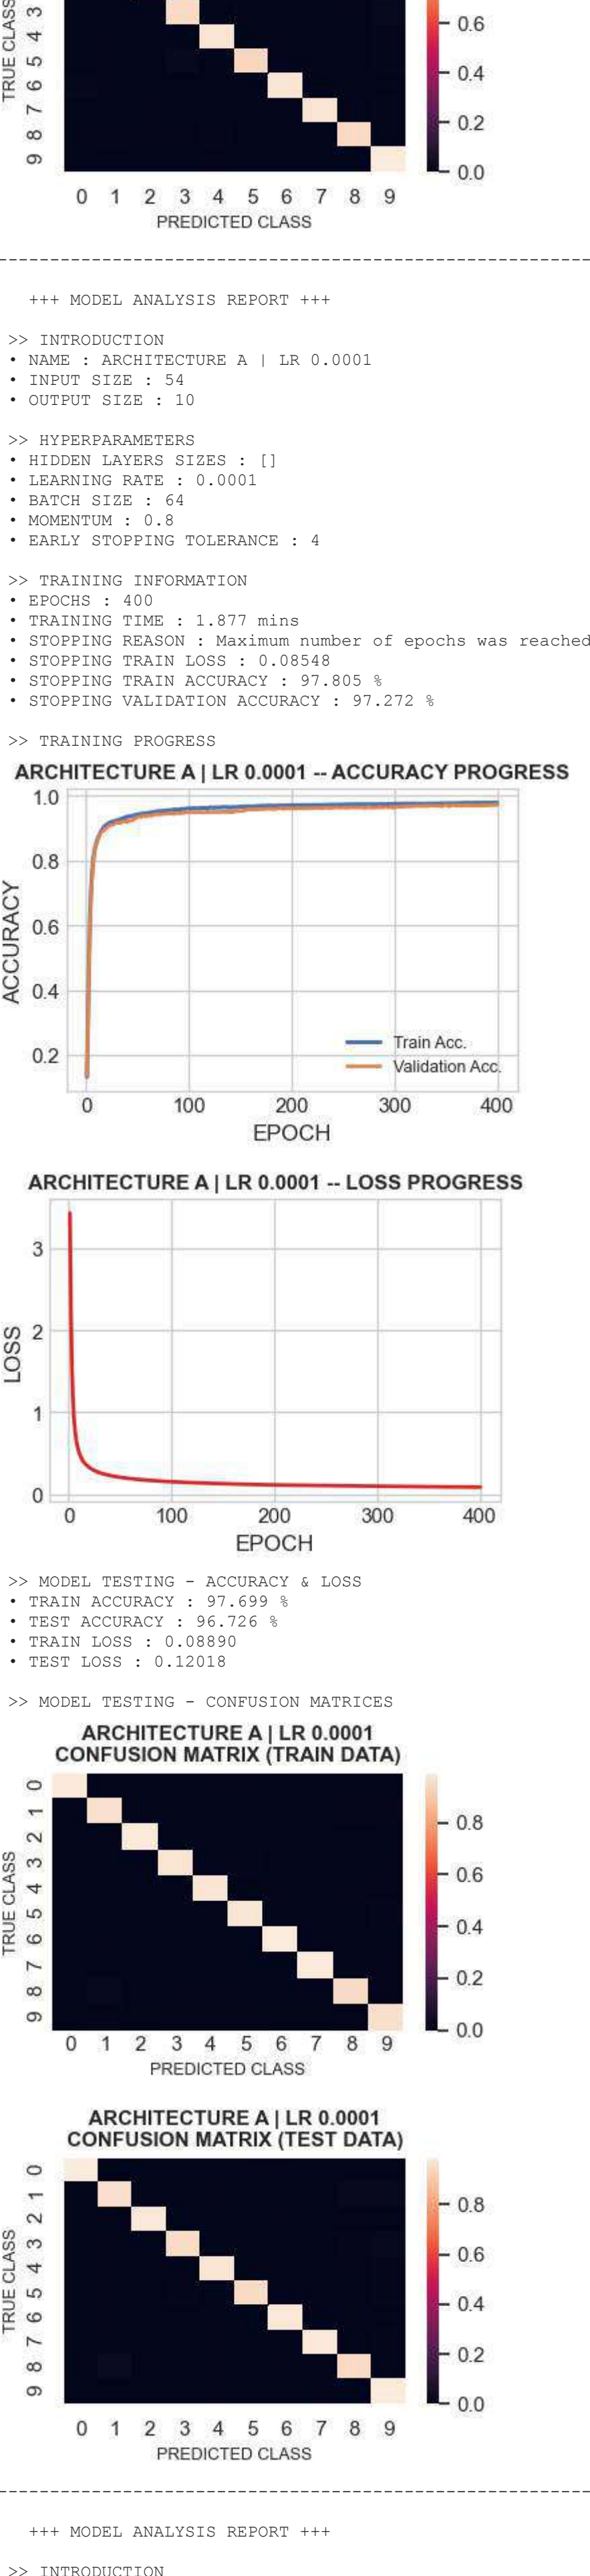

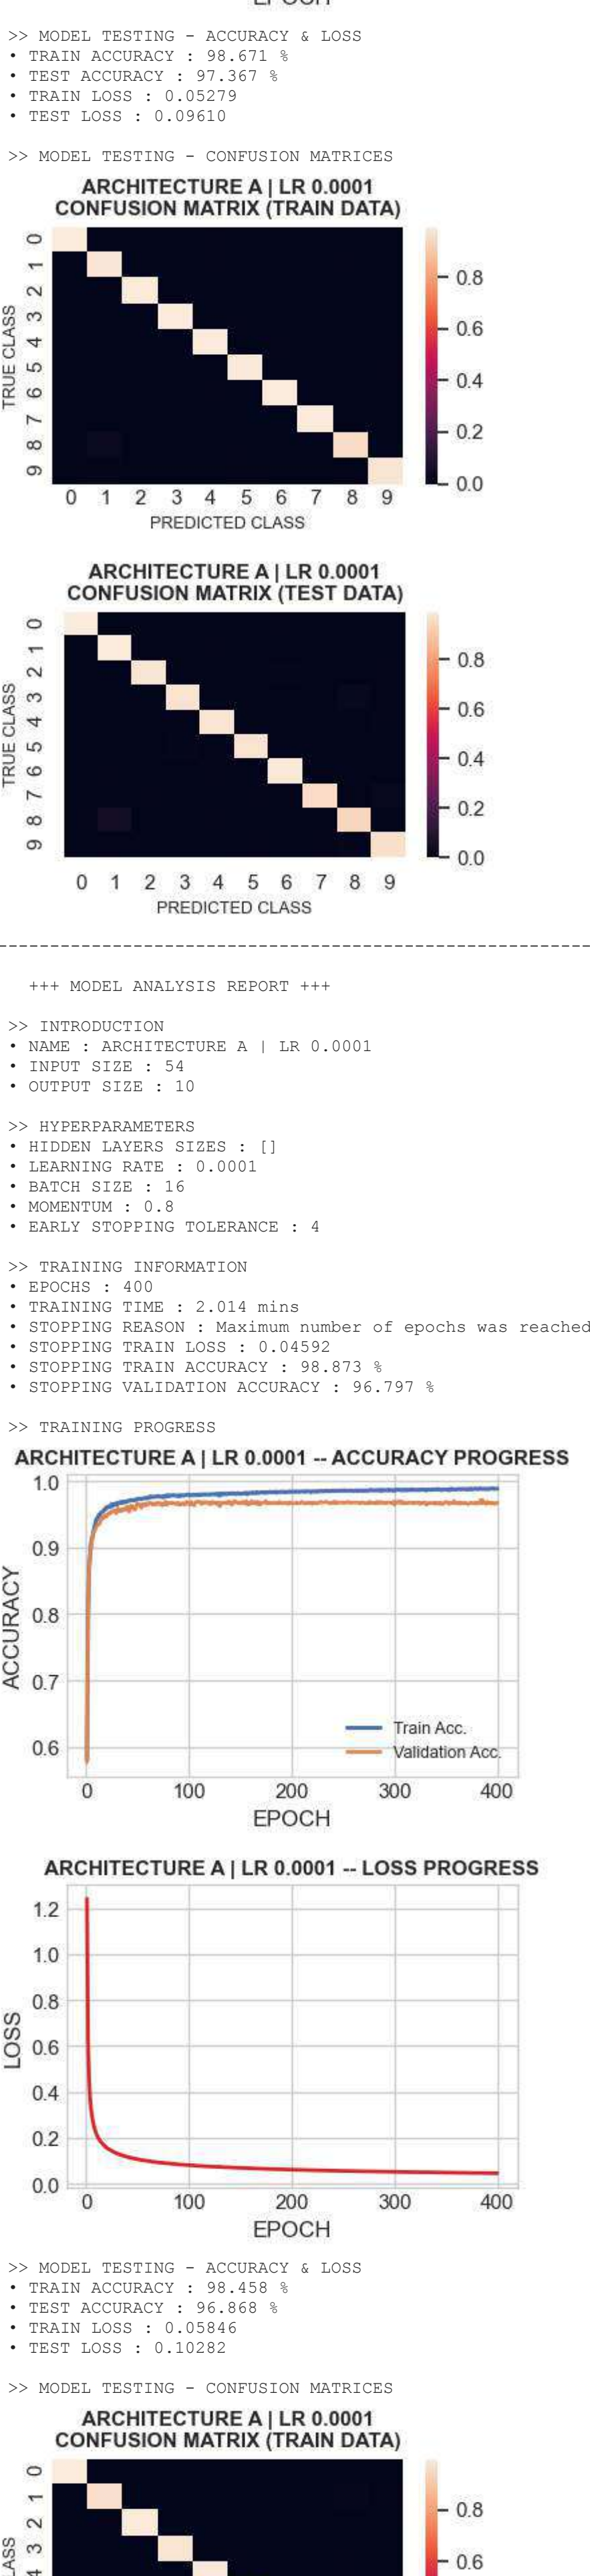

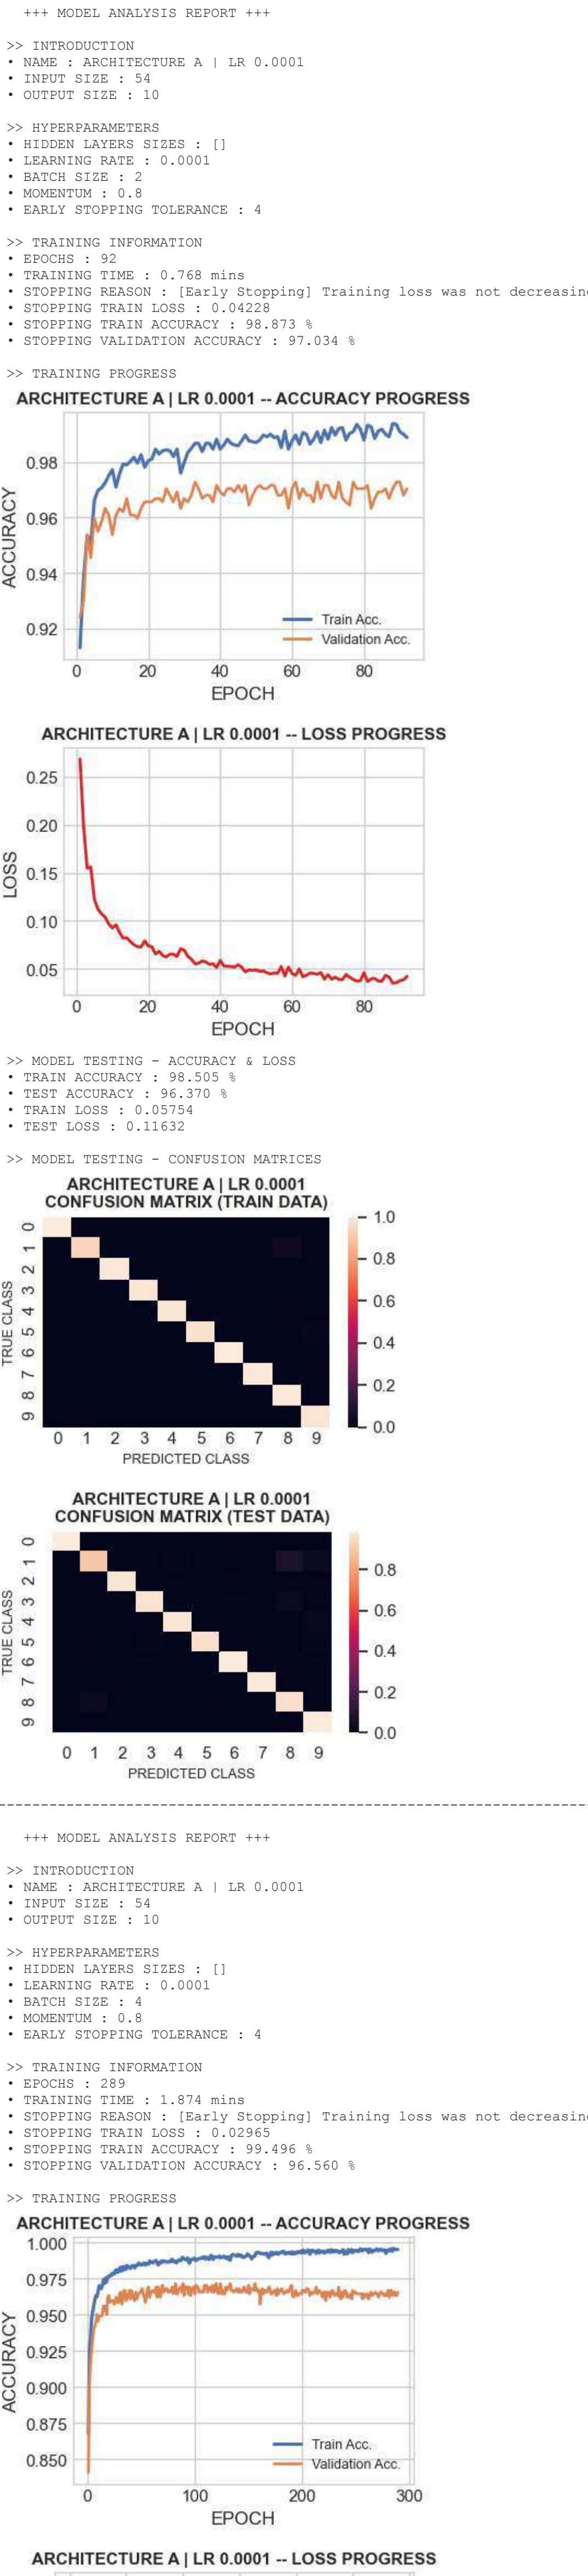
- NAME : ARCHITECTURE D | LR 0.0001
- INPUT SIZE : 54
- OUTPUT SIZE : 10

>> HYPERPARAMETERS

- HIDDEN LAYERS SIZES : [2, 3]
- LEARNING RATE : 0.0001
- BATCH SIZE : 16
- MOMENTUM : 0.8
- EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION

- EPOCHS : 400
- TRAINING TIME : 3.226 mins
- STOPPING REASON : Maximum number of epochs was reached
- STOPPING TRAIN LOSS :



```
momentums = [0.5, 0.6, 0.7, 0.8, 0.9, 1.0]
highest_test_acc_obs = 0.0
best_momentum = None
for mom in momentums:
    print(f"<div> * 50, '\n')
    model = Trainer(train_data, momentum = mom)
    test_acc = model_model_analysis_report(train_data, test_data)
    if test_acc > highest_test_acc_obs:
        highest_test_acc_obs = test_acc
        best_momentum = mom
print(f"<div> * 50, '\n')
```



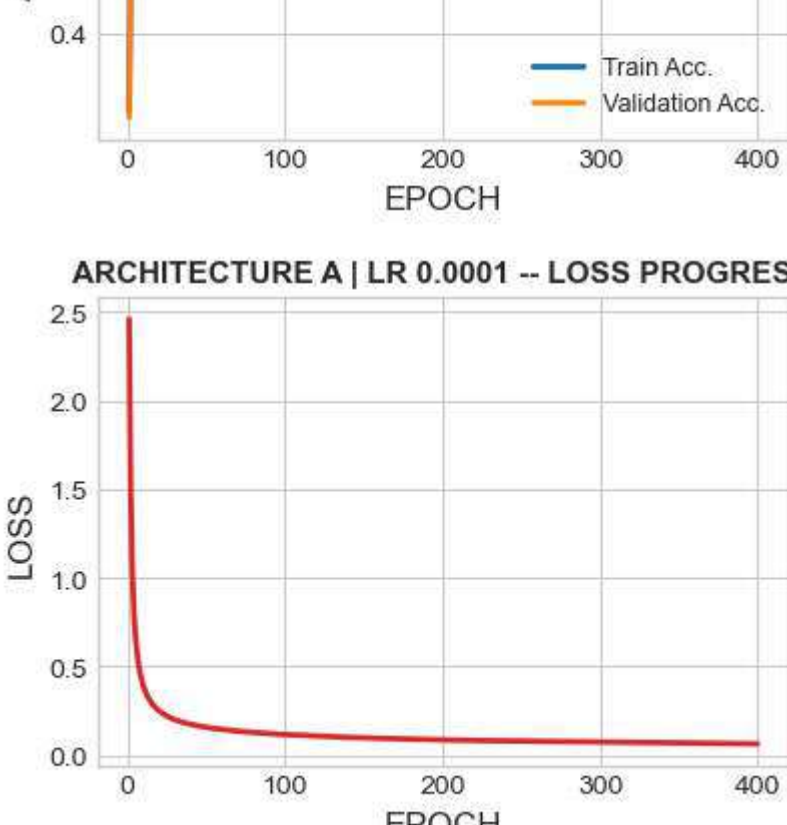
```
+++ MODEL ANALYSIS REPORT +++

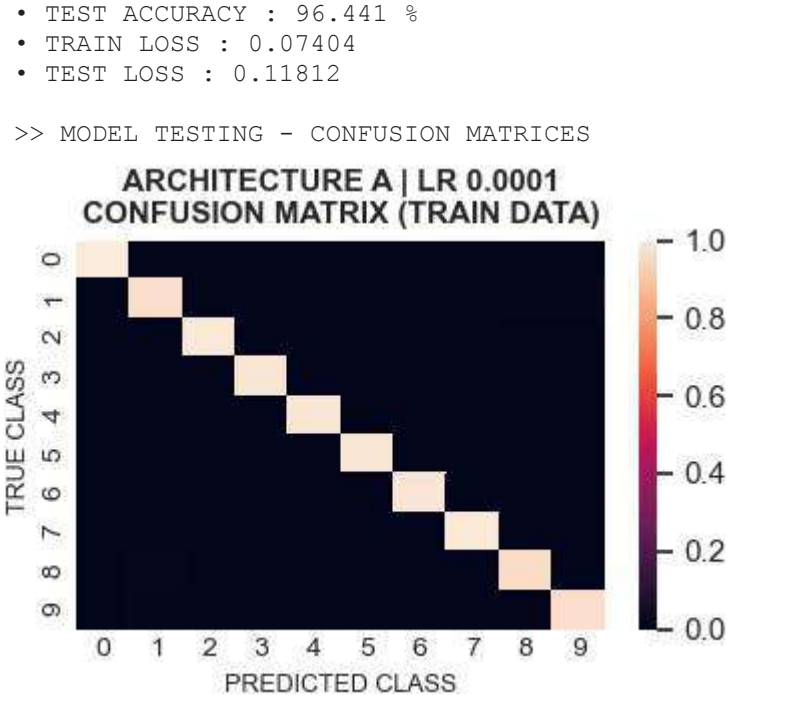
>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.5
• EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 1.991 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 0.06430
• STOPPING TRAIN ACCURACY : 98.488 %
• STOPPING VALIDATION ACCURACY : 96.560 %

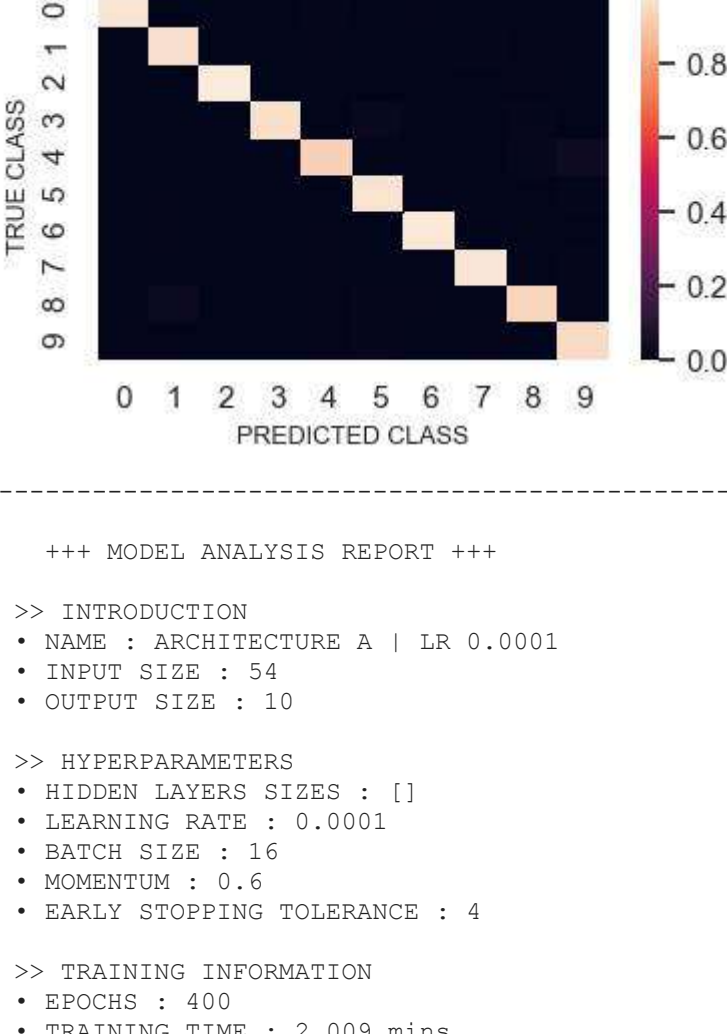
>> TRAINING PROGRESS

ARCHITECTURE A | LR 0.0001 -- ACCURACY PROGRESS


ARCHITECTURE A | LR 0.0001 -- LOSS PROGRESS


>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 98.102 %
• TEST ACCURACY : 96.441 %
• TRAIN LOSS : 0.07404
• TEST LOSS : 0.11812

>> MODEL TESTING - CONFUSION MATRICES
```



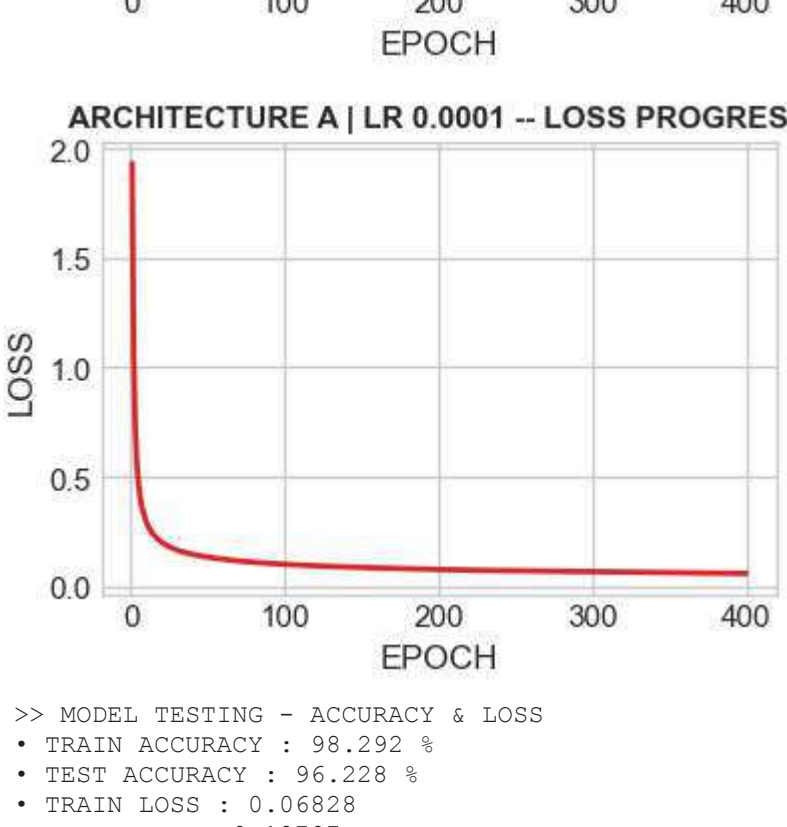
```
+++ MODEL ANALYSIS REPORT +++

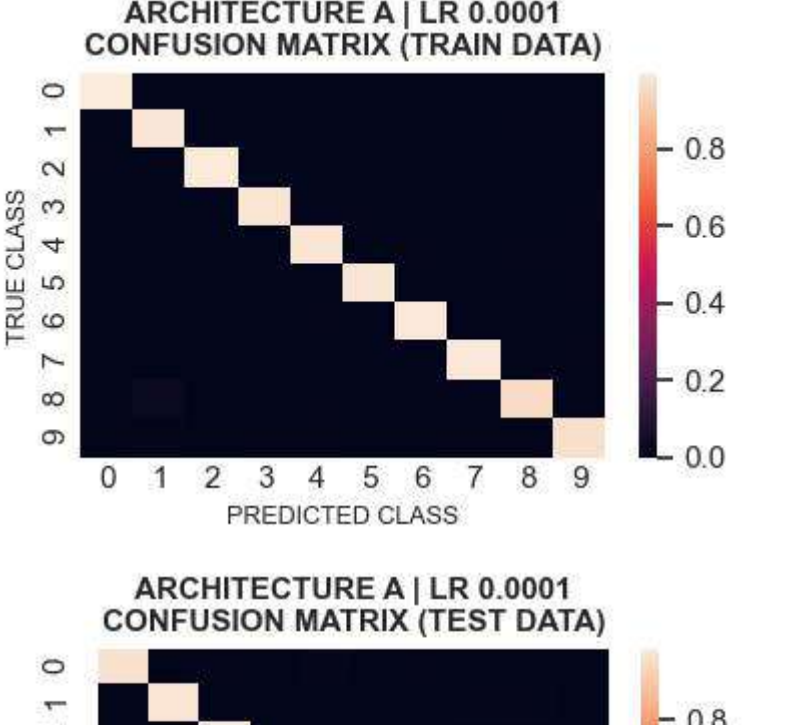
>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.6
• EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 2.009 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 0.05712
• STOPPING TRAIN ACCURACY : 98.665 %
• STOPPING VALIDATION ACCURACY : 96.797 %


>> TRAINING PROGRESS

ARCHITECTURE A | LR 0.0001 -- ACCURACY PROGRESS


ARCHITECTURE A | LR 0.0001 -- LOSS PROGRESS


>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 98.232 %
• TEST ACCURACY : 96.228 %
• TRAIN LOSS : 0.06828
• TEST LOSS : 0.12797

>> MODEL TESTING - CONFUSION MATRICES
```



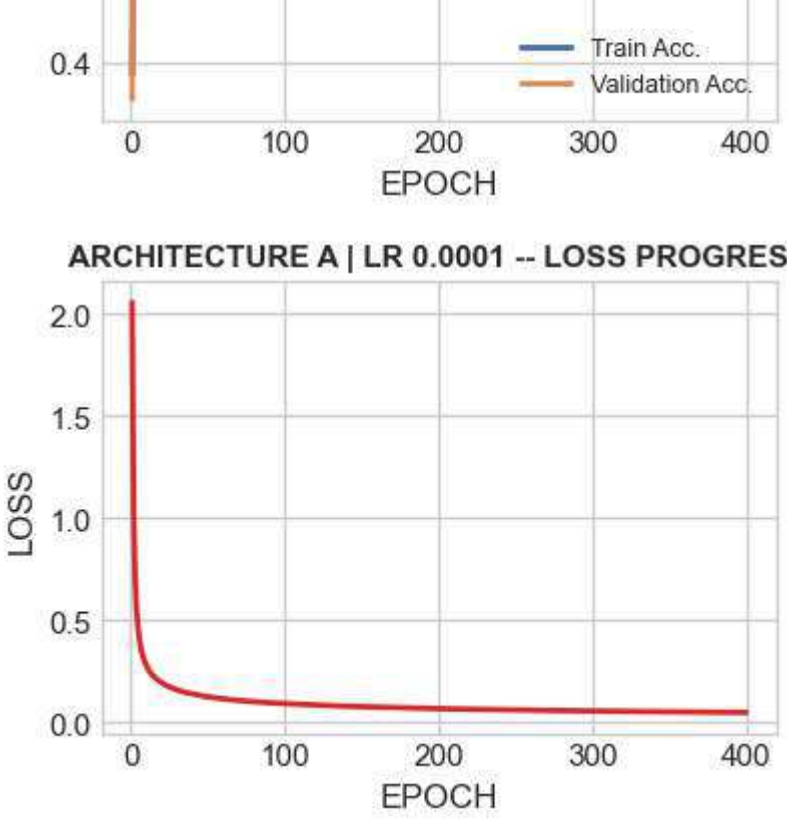
```
+++ MODEL ANALYSIS REPORT +++

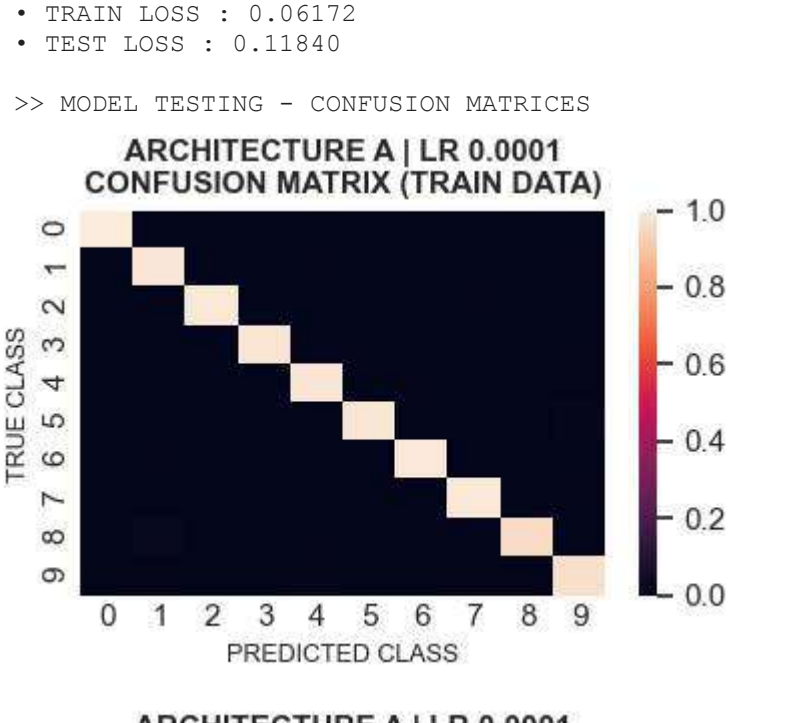
>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.7
• EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 2.006 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 0.04932
• STOPPING TRAIN ACCURACY : 98.814 %
• STOPPING VALIDATION ACCURACY : 97.034 %

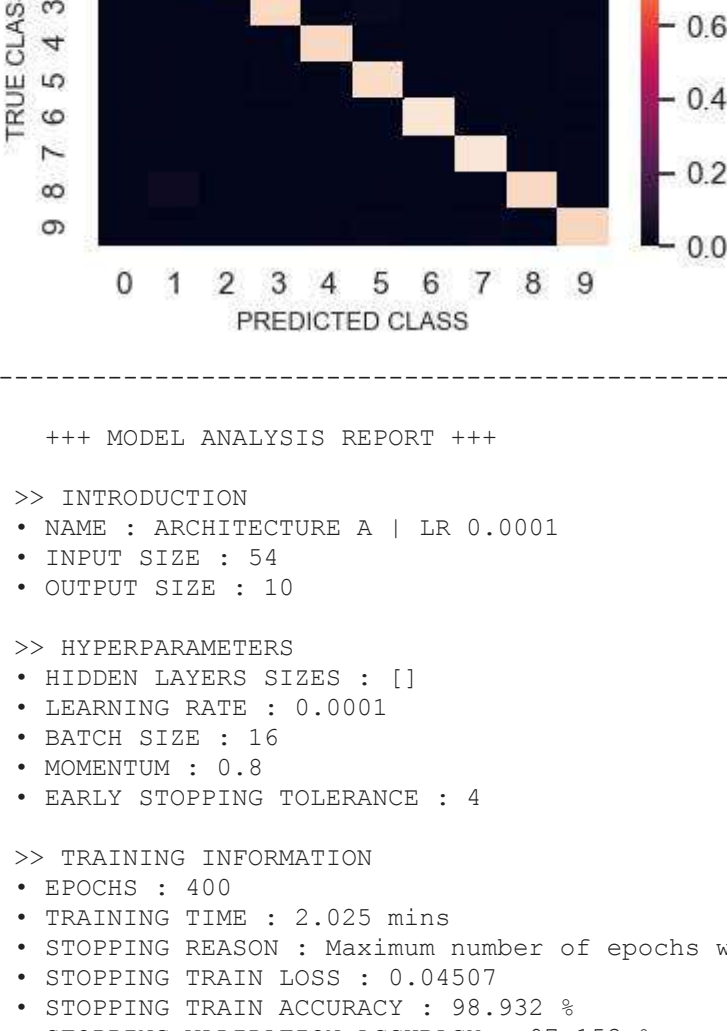
>> TRAINING PROGRESS

ARCHITECTURE A | LR 0.0001 -- ACCURACY PROGRESS


ARCHITECTURE A | LR 0.0001 -- LOSS PROGRESS


>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 98.458 %
• TEST ACCURACY : 96.726 %
• TRAIN LOSS : 0.06172
• TEST LOSS : 0.11840

>> MODEL TESTING - CONFUSION MATRICES
```



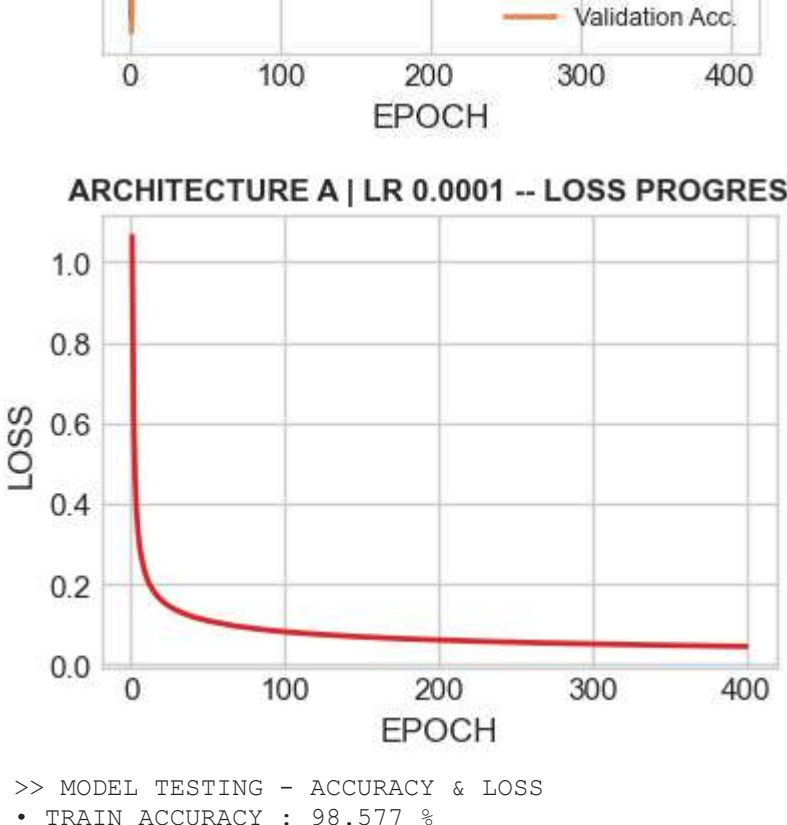
```
+++ MODEL ANALYSIS REPORT +++

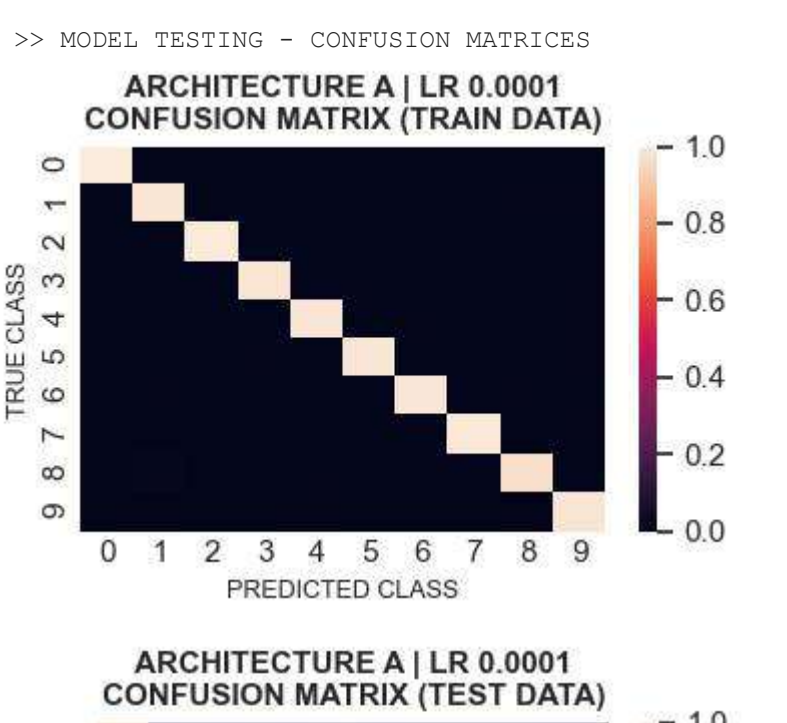
>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 2.025 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 0.04507
• STOPPING TRAIN ACCURACY : 98.932 %
• STOPPING VALIDATION ACCURACY : 97.153 %

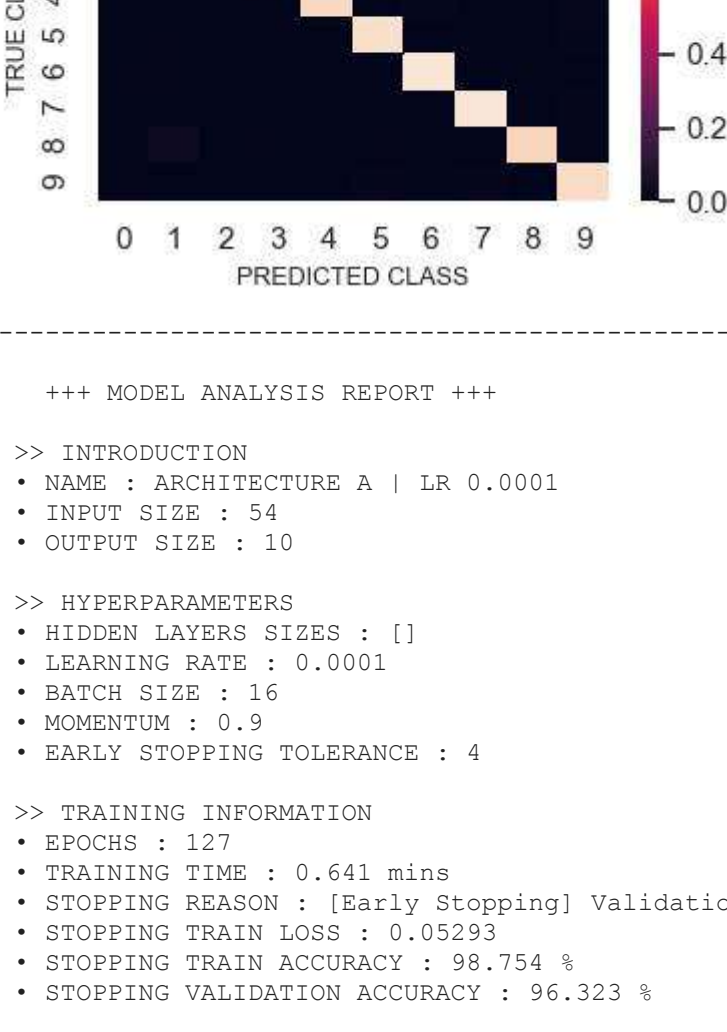
>> TRAINING PROGRESS

ARCHITECTURE A | LR 0.0001 -- ACCURACY PROGRESS


ARCHITECTURE A | LR 0.0001 -- LOSS PROGRESS


>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 98.577 %
• TEST ACCURACY : 96.940 %
• TRAIN LOSS : 0.05301
• TEST LOSS : 0.11243

>> MODEL TESTING - CONFUSION MATRICES
```



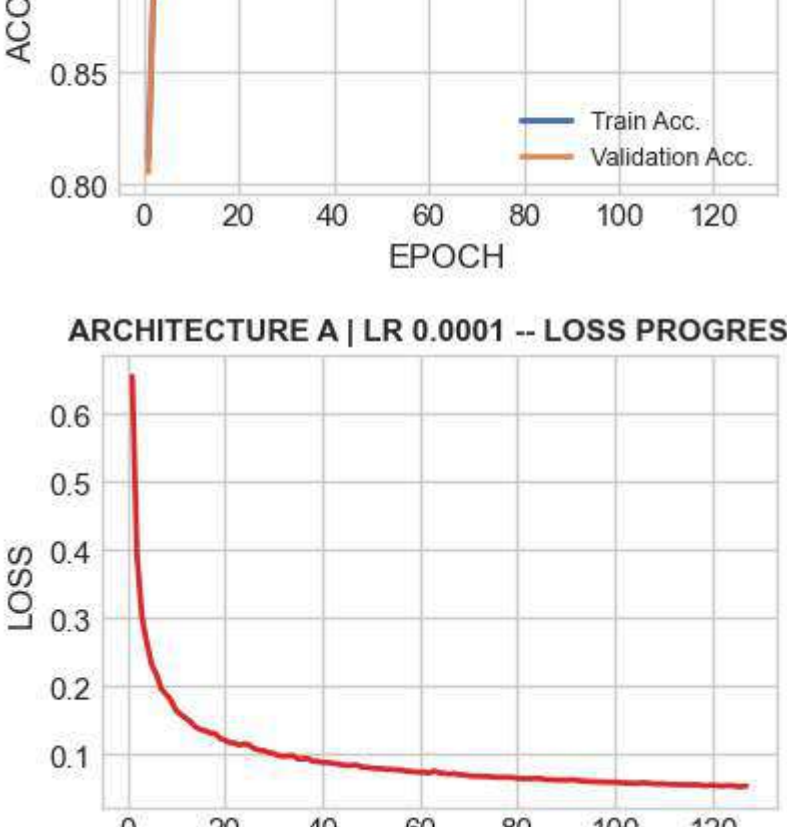
```
+++ MODEL ANALYSIS REPORT +++

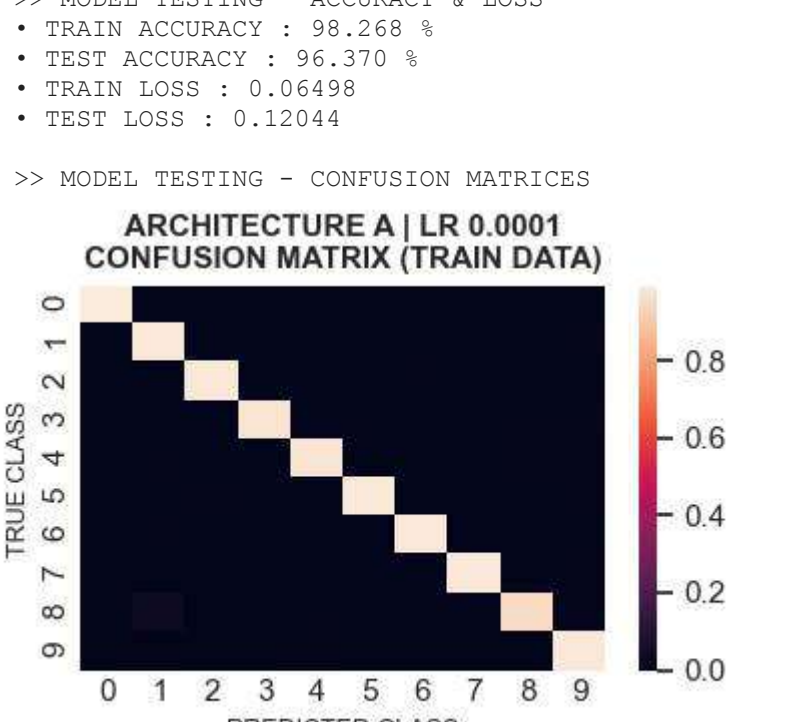
>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 1.0
• EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION
• EPOCHS : 127
• TRAINING TIME : 0.641 mins
• STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
• STOPPING TRAIN LOSS : 0.05293
• STOPPING TRAIN ACCURACY : 98.754 %
• STOPPING VALIDATION ACCURACY : 96.323 %


>> TRAINING PROGRESS

ARCHITECTURE A | LR 0.0001 -- ACCURACY PROGRESS


ARCHITECTURE A | LR 0.0001 -- LOSS PROGRESS


>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 98.268 %
• TEST ACCURACY : 96.370 %
• TRAIN LOSS : 0.06498
• TEST LOSS : 0.12044

>> MODEL TESTING - CONFUSION MATRICES
```



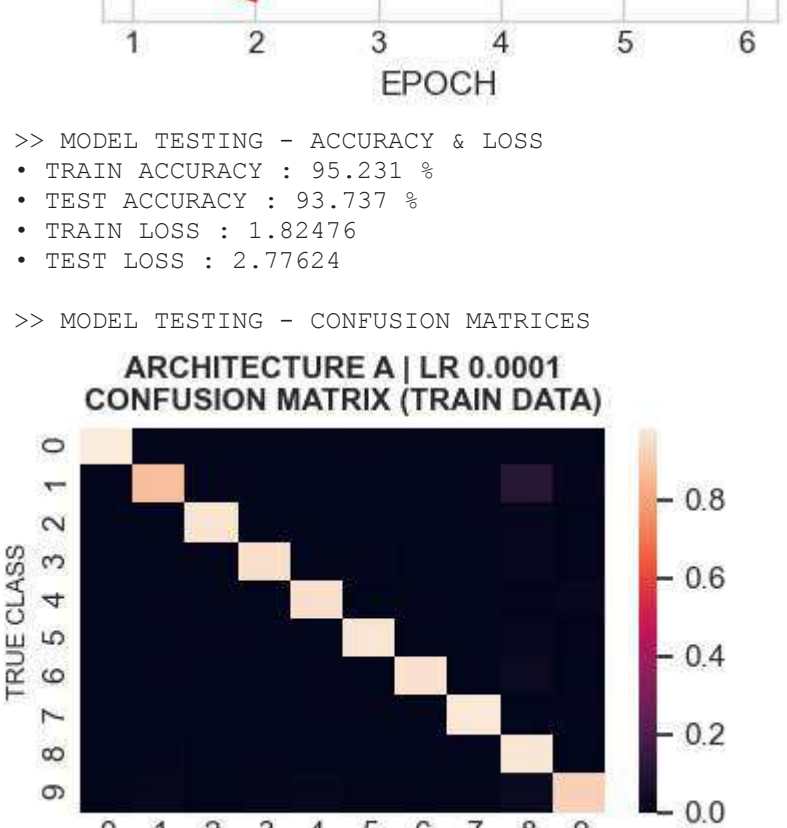
```
+++ MODEL ANALYSIS REPORT +++

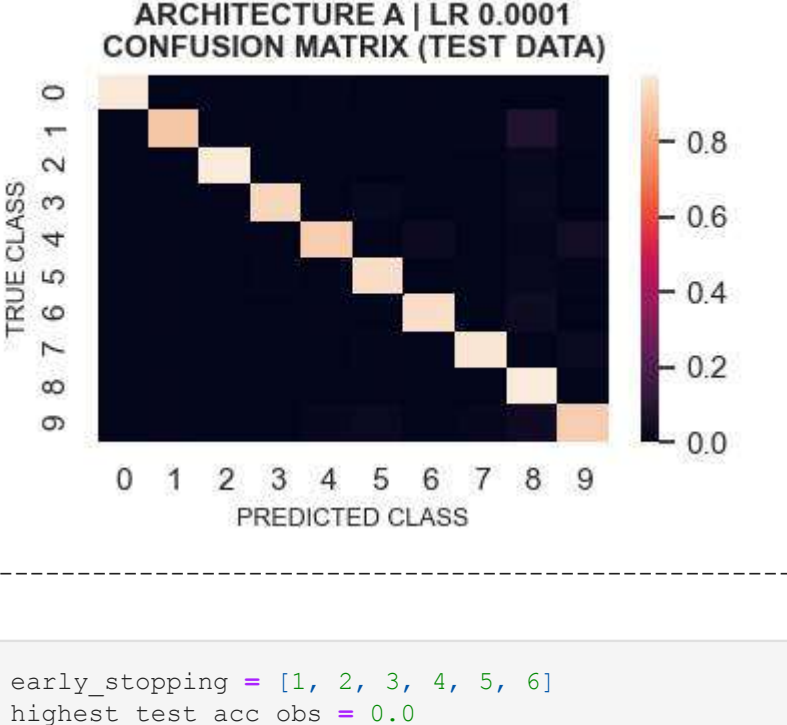
>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 54
• OUTPUT SIZE : 10

>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : []
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 1.0
• EARLY STOPPING TOLERANCE : 4

>> TRAINING INFORMATION
• EPOCHS : 6
• TRAINING TIME : 0.030 mins
• STOPPING REASON : [Early Stopping] Training loss was not decreasing
• STOPPING TRAIN LOSS : 1.73402
• STOPPING TRAIN ACCURACY : 95.344 %
• STOPPING VALIDATION ACCURACY : 94.781 %


>> TRAINING PROGRESS

ARCHITECTURE A | LR 0.0001 -- ACCURACY PROGRESS


ARCHITECTURE A | LR 0.0001 -- LOSS PROGRESS



>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 95.231 %
• TEST ACCURACY : 93.737 %
• TRAIN LOSS : 1.82476
• TEST LOSS : 2.77624

>> MODEL TESTING - CONFUSION MATRICES
```



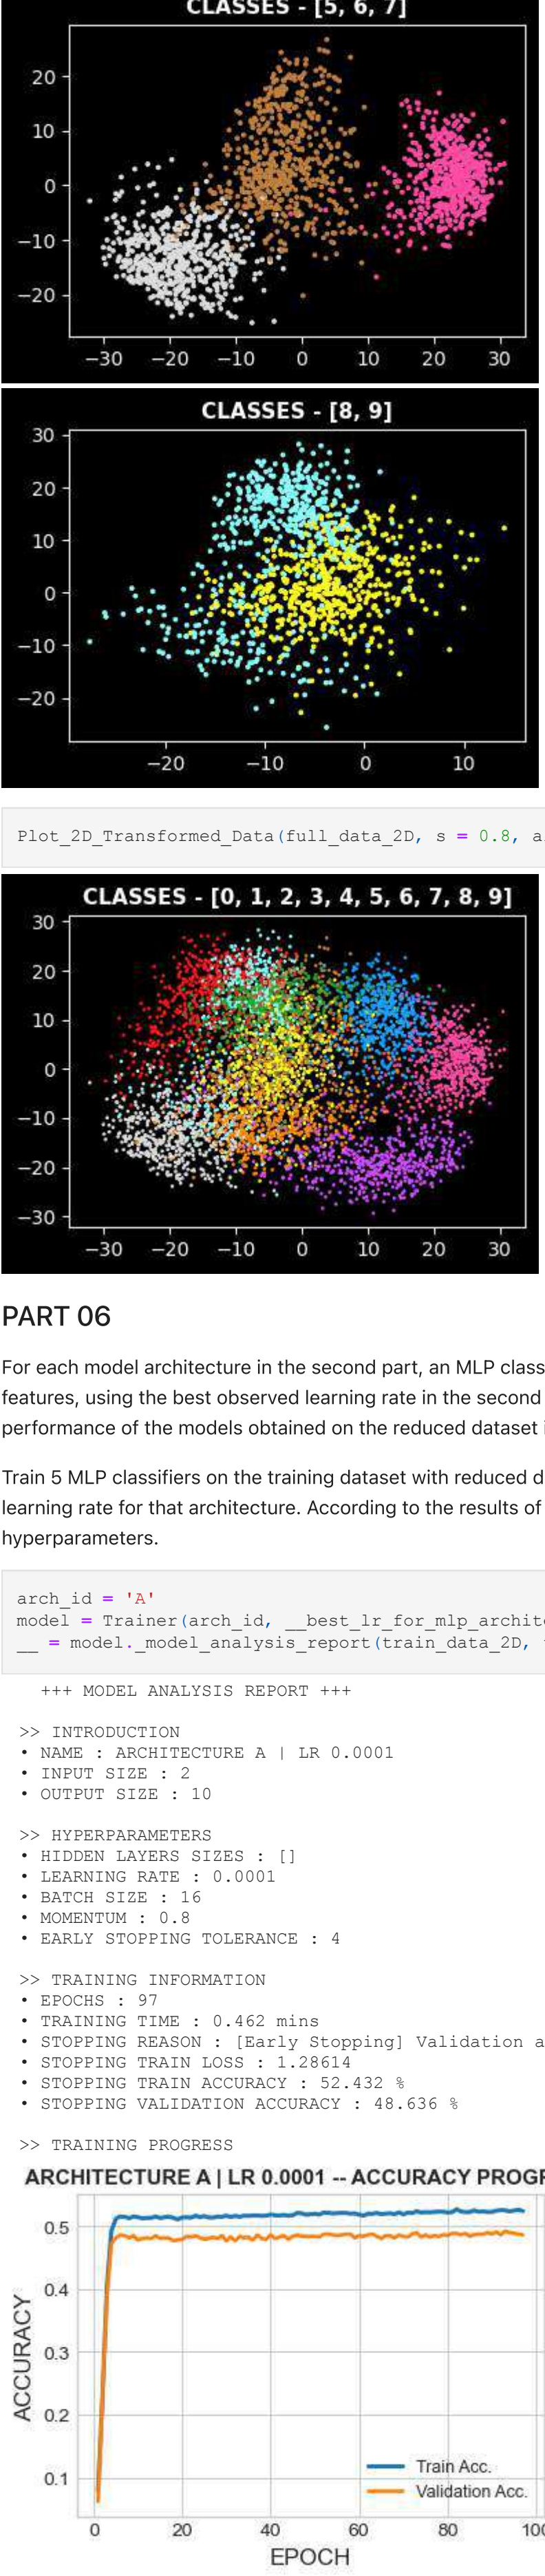
In [28]:

```
early_stopping = [1, 2, 3, 4, 5, 6]
highest_test_acc_obs = 0.0
best_early_stopping_tol = None
for tol in early_stopping:
    print('--' * 50, '\n')
    model = Trainer( best_architecture_found , best_learning_rate_found ,
                      train_data, early_stopping_tol = tol)
    test_acc = model._model_analysis_report(train_data, test_data)
    if test_acc > highest_test_acc_obs:
        highest_test_acc_obs = test_acc
        best_early_stopping_tol = tol
    print('--' * 50, '\n')
```

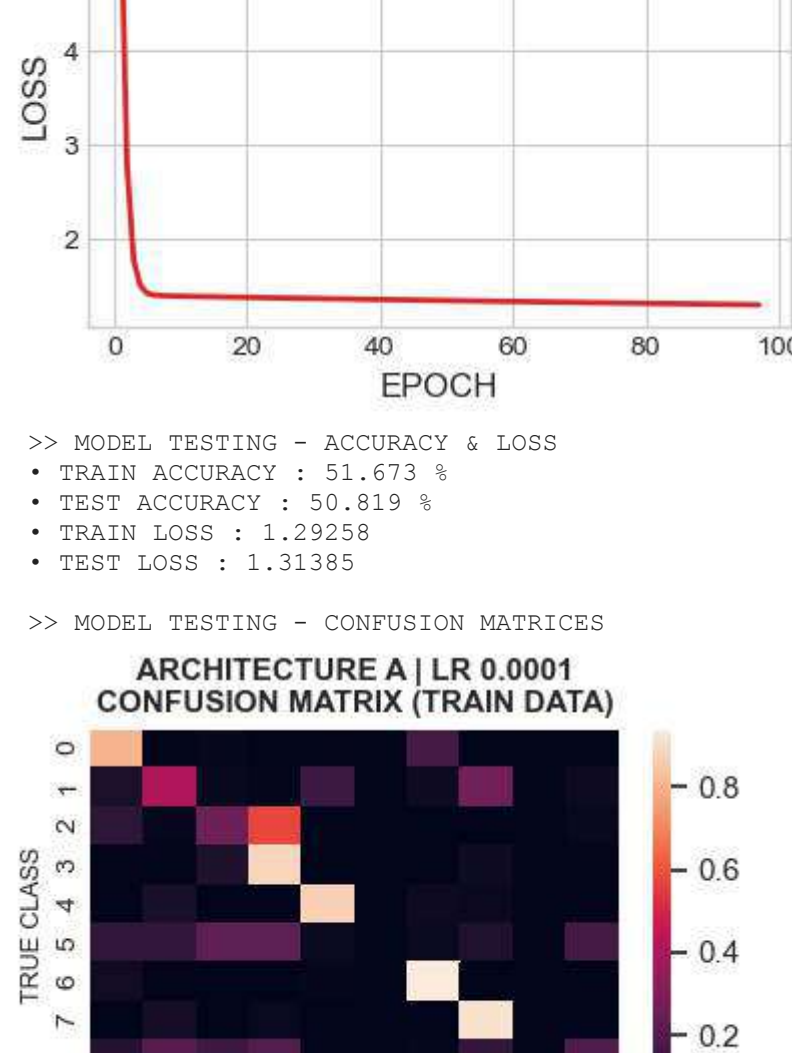
In [35]:

```
Plot_2D_Transformed_Data(full_data_2D, [0, 1, 4])
Plot_2D_Transformed_Data(full_data_2D, [2, 3])
Plot_2D_Transformed_Data(full_data_2D, [5, 6, 7])
Plot_2D_Transformed_Data(full_data_2D, [8, 9])
```



In [36]:

```
Plot_2D_Transformed_Data(full_data_2D, s = 0.8, alpha = 1)
```



PART 06

For each model architecture in the second part, an MLP classifier is trained on the dataset with reduced dimensionality of the features using the best observed learning rate in the second part as the learning rate for the respective model architecture. The performance of the models obtained on the reduced dataset is compared with the earlier performance.

Train 5 MLP classifiers on the training dataset with reduced dimensions, each with a different model architecture and the best learning rate for that architecture. According to the results of the second part, the 5 models should have the following hyperparameters.

In [39]:

```
arch_id = 'B'
model = Trainer(arch_id, _best_lr_for_mlp_architectures_[arch_id], train_data_2D)
model._model_analysis_report(train_data_2D, test_data_2D)

+++ MODEL ANALYSIS REPORT +++

>>> INTRODUCTION
• NAME : ARCHITECTURE A | LR 0.0001
• INPUT SIZE : 2
• OUTPUT SIZE : 10

>>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : {}
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

>>> TRAINING INFORMATION
• EPOCHS : 97
• TRAINING TIME : 0.462 mins
• STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
• STOPPING TRAIN LOSS : 1.3252
• STOPPING TRAIN ACCURACY : 51.432 %
• STOPPING VALIDATION ACCURACY : 48.636 %

>>> TRAINING PROGRESS

ARCHITECTURE A | LR 0.0001 -- ACCURACY PROGRESS

ACCURACY
0.5
0.4
0.3
0.2
0.1
0
0 20 40 60 80 100
EPOCH
Train Acc.
Validation Acc.

ARCHITECTURE A | LR 0.0001 -- LOSS PROGRESS

LOSS
6
5
4
3
2
1
0 20 40 60 80 100
EPOCH

>>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 51.673 %
• TEST ACCURACY : 50.819 %
• TRAIN LOSS : 1.32958
• TEST LOSS : 1.31385

>>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE A | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS

ARCHITECTURE A | LR 0.0001
CONFUSION MATRIX (TEST DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS
```

In [39]:

```
arch_id = 'B'
model = Trainer(arch_id, _best_lr_for_mlp_architectures_[arch_id], train_data_2D)
model._model_analysis_report(train_data_2D, test_data_2D)

+++ MODEL ANALYSIS REPORT +++

>>> INTRODUCTION
• NAME : ARCHITECTURE B | LR 0.0001
• INPUT SIZE : 2
• OUTPUT SIZE : 10

>>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : {}
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

>>> TRAINING INFORMATION
• EPOCHS : 160
• TRAINING TIME : 0.989 mins
• STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
• STOPPING TRAIN LOSS : 1.3252
• STOPPING TRAIN ACCURACY : 51.512 %
• STOPPING VALIDATION ACCURACY : 52.550 %

>>> TRAINING PROGRESS

ARCHITECTURE B | LR 0.0001 -- ACCURACY PROGRESS

ACCURACY
0.5
0.4
0.3
0.2
0.1
0 25 50 75 100 125 150
EPOCH
Train Acc.
Validation Acc.

ARCHITECTURE B | LR 0.0001 -- LOSS PROGRESS

LOSS
4.5
4
3.5
3
2.5
2
1.5
0 25 50 75 100 125 150
EPOCH

>>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 51.720 %
• TEST ACCURACY : 51.174 %
• TRAIN LOSS : 1.32960
• TEST LOSS : 1.34624

>>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE B | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS

ARCHITECTURE B | LR 0.0001
CONFUSION MATRIX (TEST DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS
```

In [40]:

```
arch_id = 'C'
model = Trainer(arch_id, _best_lr_for_mlp_architectures_[arch_id], train_data_2D)
model._model_analysis_report(train_data_2D, test_data_2D)

+++ MODEL ANALYSIS REPORT +++

>>> INTRODUCTION
• NAME : ARCHITECTURE C | LR 0.0001
• INPUT SIZE : 2
• OUTPUT SIZE : 10

>>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : {}
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

>>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 2.496 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 1.13989
• STOPPING TRAIN ACCURACY : 58.096 %
• STOPPING VALIDATION ACCURACY : 58.126 %

>>> TRAINING PROGRESS

ARCHITECTURE C | LR 0.0001 -- ACCURACY PROGRESS

ACCURACY
0.55
0.5
0.45
0.4
0.35
0.3
0.25
0.2
0.15
0.1
0 100 200 300 400
EPOCH
Train Acc.
Validation Acc.

ARCHITECTURE C | LR 0.0001 -- LOSS PROGRESS

LOSS
2.2
2.0
1.8
1.6
1.4
1.2
1.0
0 100 200 300 400
EPOCH

>>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 58.102 %
• TEST ACCURACY : 55.730 %
• TRAIN LOSS : 1.13433
• TEST LOSS : 1.15921

>>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE C | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS

ARCHITECTURE C | LR 0.0001
CONFUSION MATRIX (TEST DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS
```

In [41]:

```
arch_id = 'D'
model = Trainer(arch_id, _best_lr_for_mlp_architectures_[arch_id], train_data_2D)
model._model_analysis_report(train_data_2D, test_data_2D)

+++ MODEL ANALYSIS REPORT +++

>>> INTRODUCTION
• NAME : ARCHITECTURE D | LR 0.01
• INPUT SIZE : 2
• OUTPUT SIZE : 10

>>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : {}
• LEARNING RATE : 0.01
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

>>> TRAINING INFORMATION
• EPOCHS : 125
• TRAINING TIME : 0.962 mins
• STOPPING REASON : [Early Stopping] Validation accuracy was not increasing
• STOPPING TRAIN LOSS : 1.06826
• STOPPING TRAIN ACCURACY : 57.117 %
• STOPPING VALIDATION ACCURACY : 63.582 %

>>> TRAINING PROGRESS

ARCHITECTURE D | LR 0.01 -- ACCURACY PROGRESS

ACCURACY
0.65
0.6
0.55
0.5
0.45
0.4
0.35
0.3
0.25
0.2
0.15
0.1
0 25 50 75 100 125
EPOCH
Train Acc.
Validation Acc.

ARCHITECTURE D | LR 0.01 -- LOSS PROGRESS

LOSS
1.4
1.3
1.2
1.1
1.0
0 20 40 60 80 100 120
EPOCH

>>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 59.644 %
• TEST ACCURACY : 57.224 %
• TRAIN LOSS : 1.05428
• TEST LOSS : 1.08120

>>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE D | LR 0.01
CONFUSION MATRIX (TRAIN DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS

ARCHITECTURE D | LR 0.01
CONFUSION MATRIX (TEST DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS
```

In [42]:

```
arch_id = 'E'
model = Trainer(arch_id, _best_lr_for_mlp_architectures_[arch_id], train_data_2D)
model._model_analysis_report(train_data_2D, test_data_2D)

+++ MODEL ANALYSIS REPORT +++

>>> INTRODUCTION
• NAME : ARCHITECTURE E | LR 0.0001
• INPUT SIZE : 2
• OUTPUT SIZE : 10

>>> HYPERPARAMETERS
• HIDDEN LAYERS SIZES : {}
• LEARNING RATE : 0.0001
• BATCH SIZE : 16
• MOMENTUM : 0.8
• EARLY STOPPING TOLERANCE : 4

>>> TRAINING INFORMATION
• EPOCHS : 400
• TRAINING TIME : 3.066 mins
• STOPPING REASON : Maximum number of epochs was reached
• STOPPING TRAIN LOSS : 1.14320
• STOPPING TRAIN ACCURACY : 57.117 %
• STOPPING VALIDATION ACCURACY : 56.465 %

>>> TRAINING PROGRESS

ARCHITECTURE E | LR 0.0001 -- ACCURACY PROGRESS

ACCURACY
0.5
0.4
0.3
0.2
0.1
0 100 200 300 400
EPOCH
Train Acc.
Validation Acc.

ARCHITECTURE E | LR 0.0001 -- LOSS PROGRESS

LOSS
4
3
2
1
0 100 200 300 400
EPOCH

>>> MODEL TESTING - ACCURACY & LOSS
• TRAIN ACCURACY : 56.987 %
• TEST ACCURACY : 55.943 %
• TRAIN LOSS : 1.14874
• TEST LOSS : 1.17094

>>> MODEL TESTING - CONFUSION MATRICES

ARCHITECTURE E | LR 0.0001
CONFUSION MATRIX (TRAIN DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS

ARCHITECTURE E | LR 0.0001
CONFUSION MATRIX (TEST DATA)

TRUE CLASS
9 8 7 6 5 4 3 2 1 0
0 1 2 3 4 5 6 7 8 9
PREDICTED CLASS
```